

# **Less is More: Recursive Reasoning with Tiny Networks**

**Alexia Jolicoeur-Martineau**  
**Samsung SAIL Montréal**  
**[arXiv:2510.04871](https://arxiv.org/abs/2510.04871)**

**Xi-Wei Pan Oct.15, 2025**

# TL; DR

- A single tiny network **recursively** refines its latent reasoning state to improve answers **step by step**—achieving strong generalization through thinking depth rather than network depth.
- Why recursion helps so much remains to be explained (overfitting?).
- A supervised learning method rather than generative model.

# Contents

- Background
- Deep Equilibrium Model
- Hierarchical Reasoning Model
- Tiny Recursive Model

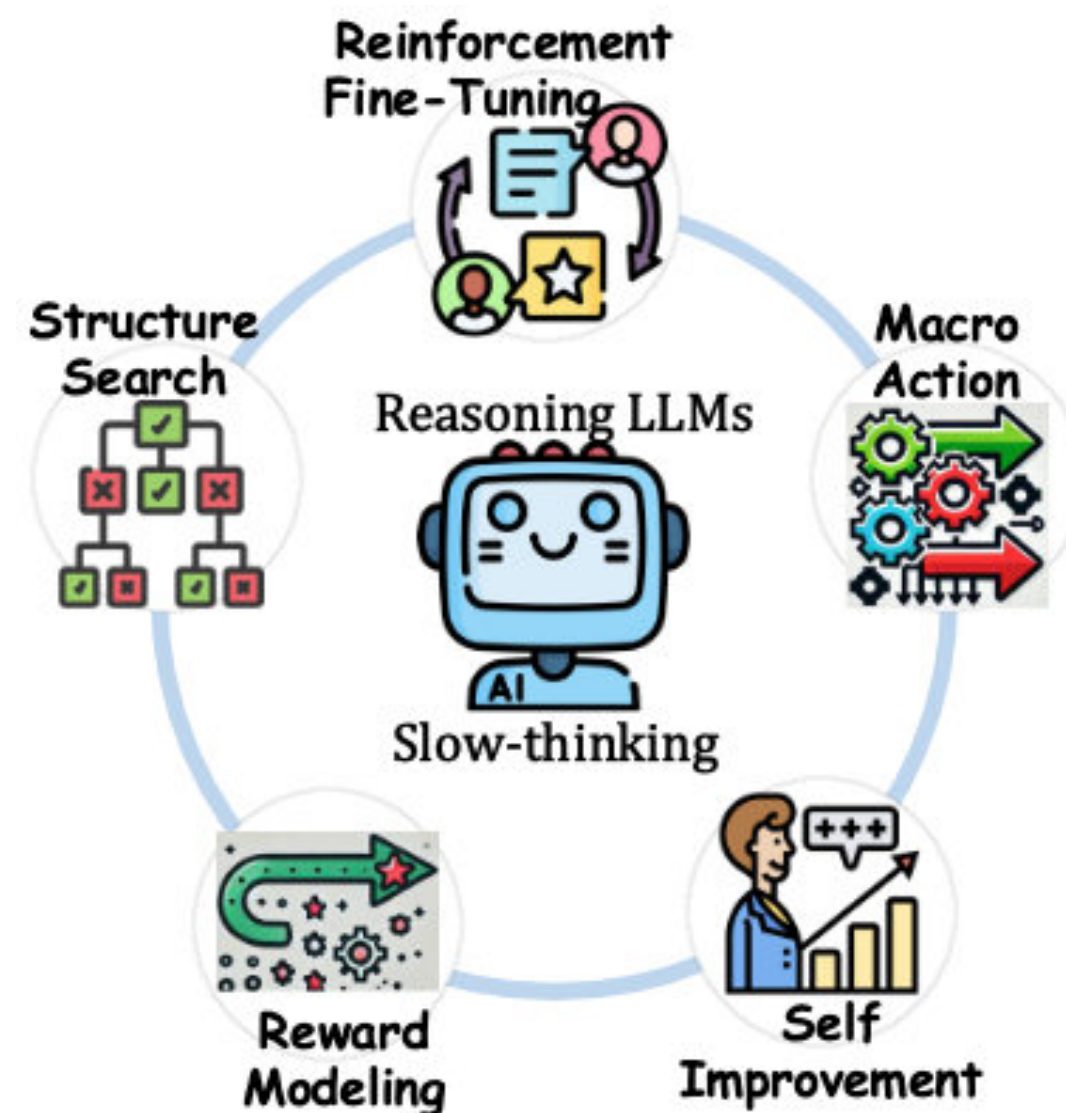
# System 1 & System 2 Thinking

- Two different modes of cognitive processing
- **System 1 is fast, automatic, intuitive, and emotional.**  
effortlessly and quickly,  
guides our daily decisions, judgments, and impressions.  
eg. Feed-forward Neural Network, Transformer?
- **System 2 is slow, deliberate, and analytical.**  
It is activated when we need to perform complex computations.  
eg. expert-level coding, competitive math, and PhD-level science questions.  
LLM: Not Turing-complete, hard to do algorithmic reasoning.

LLM + ...

# From Next-Token Prediction to Reasoning

- Traditional next-token prediction model:  $p(x_t | \mathbf{x}_{<t}) \leftarrow$  No explicit reasoning
- **Reasoning:** devising and executing complex goal oriented action sequences
  - The **Chain of Thought (CoT)** asks model to "think step by step".
  - The **Test-Time Compute (TTC)** lets model "think more" before answering.



The reasoning abilities in LLMs are typically developed through **alignment**, including

- supervised fine-tuning (SFT),
- reinforcement learning (RL).

# Deep Equilibrium Model

## Motivation - Beyond Layer Stacking

- Deep neural networks (like Transformers) rely on stacking layers for better performance.

$$z^i = f_{\theta}^{i-1}(z^{i-1}, \tilde{x}), \text{ where } i = 1, 2, \dots, L$$

- But increasing depth  $\rightarrow$  vanishing gradients, huge memory cost, slow inference.

$$f_{\theta}^i = f_{\theta}, \forall i$$

- Q: If the same transformation is applied at each layer of a deep network, what is the limit of this process?

# Deep Equilibrium Model

## Mathematical Formulation

- Instead of explicitly computing multiple layers, a DEQ finds a **fixed point**:

$$z^* = f_{\theta}(z^*, \tilde{x})$$

- Here,  $z^*$  is the representation **at equilibrium**, meaning applying further does not change  $z^*$ .
- Unlike a conventional network where the output is the activations from the  $L^{\text{th}}$  layer, **the output of a DEQ is the equilibrium point itself.**

# Deep Equilibrium Model

## Finding the Equilibrium: Root-Finding Methods

- Fixed-Point Iteration:  $z^{t+1} = f_{\theta}(z^t, \tilde{x})$       Slow!
- Broyden's Method (Quasi-Newton)

$$g_{\theta}(z, \tilde{x}) := f_{\theta}(z, \tilde{x}) - z \rightarrow 0$$

$$z^{t+1} = z^t - J_g^{-1} g(z^t)$$

- Any root-find method is fine.



# Deep Equilibrium Model

## Training DEQs with Implicit Differentiation

- Unlike traditional deep networks that require storing activations for back-propagation, DEQs train using **implicit differentiation**.

- Loss function:  $\mathcal{L} = \mathcal{L}(z^*, y)$   $\frac{d\mathcal{L}}{d\theta} = \frac{\partial \mathcal{L}}{\partial z^*} \frac{dz^*}{d\theta}$

- **Implicit Function Theorem**  $g_\theta(z, \tilde{x}) := f_\theta(z, \tilde{x}) - z$

$$\frac{dg}{d\theta} = \frac{\partial f}{\partial \theta} + \frac{\partial f}{\partial z} \frac{dz^*}{d\theta} - \frac{dz^*}{d\theta} = 0$$

$$(I - J_f) \frac{dz^*}{d\theta} = \frac{\partial f}{\partial \theta}$$

# Hierarchical Reasoning Model

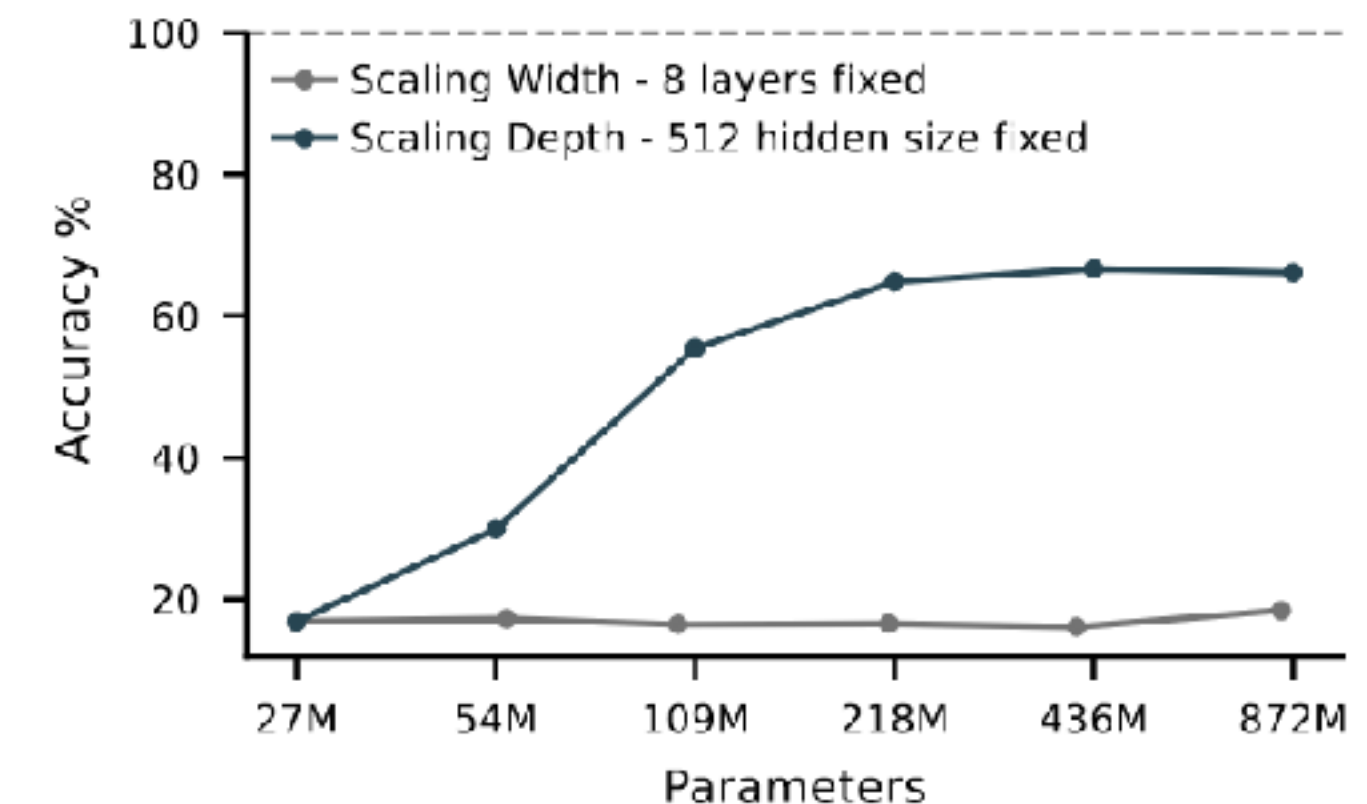

## Motivation

**Reasoning:** devising and executing complex goal oriented action sequences

**Inspiration from the brain:** signals at different frequencies + recurrence

Gap in previous work:

- CoT suffer from brittle task decomposition, high data and compute requirements
- (Transformers don't scale well)



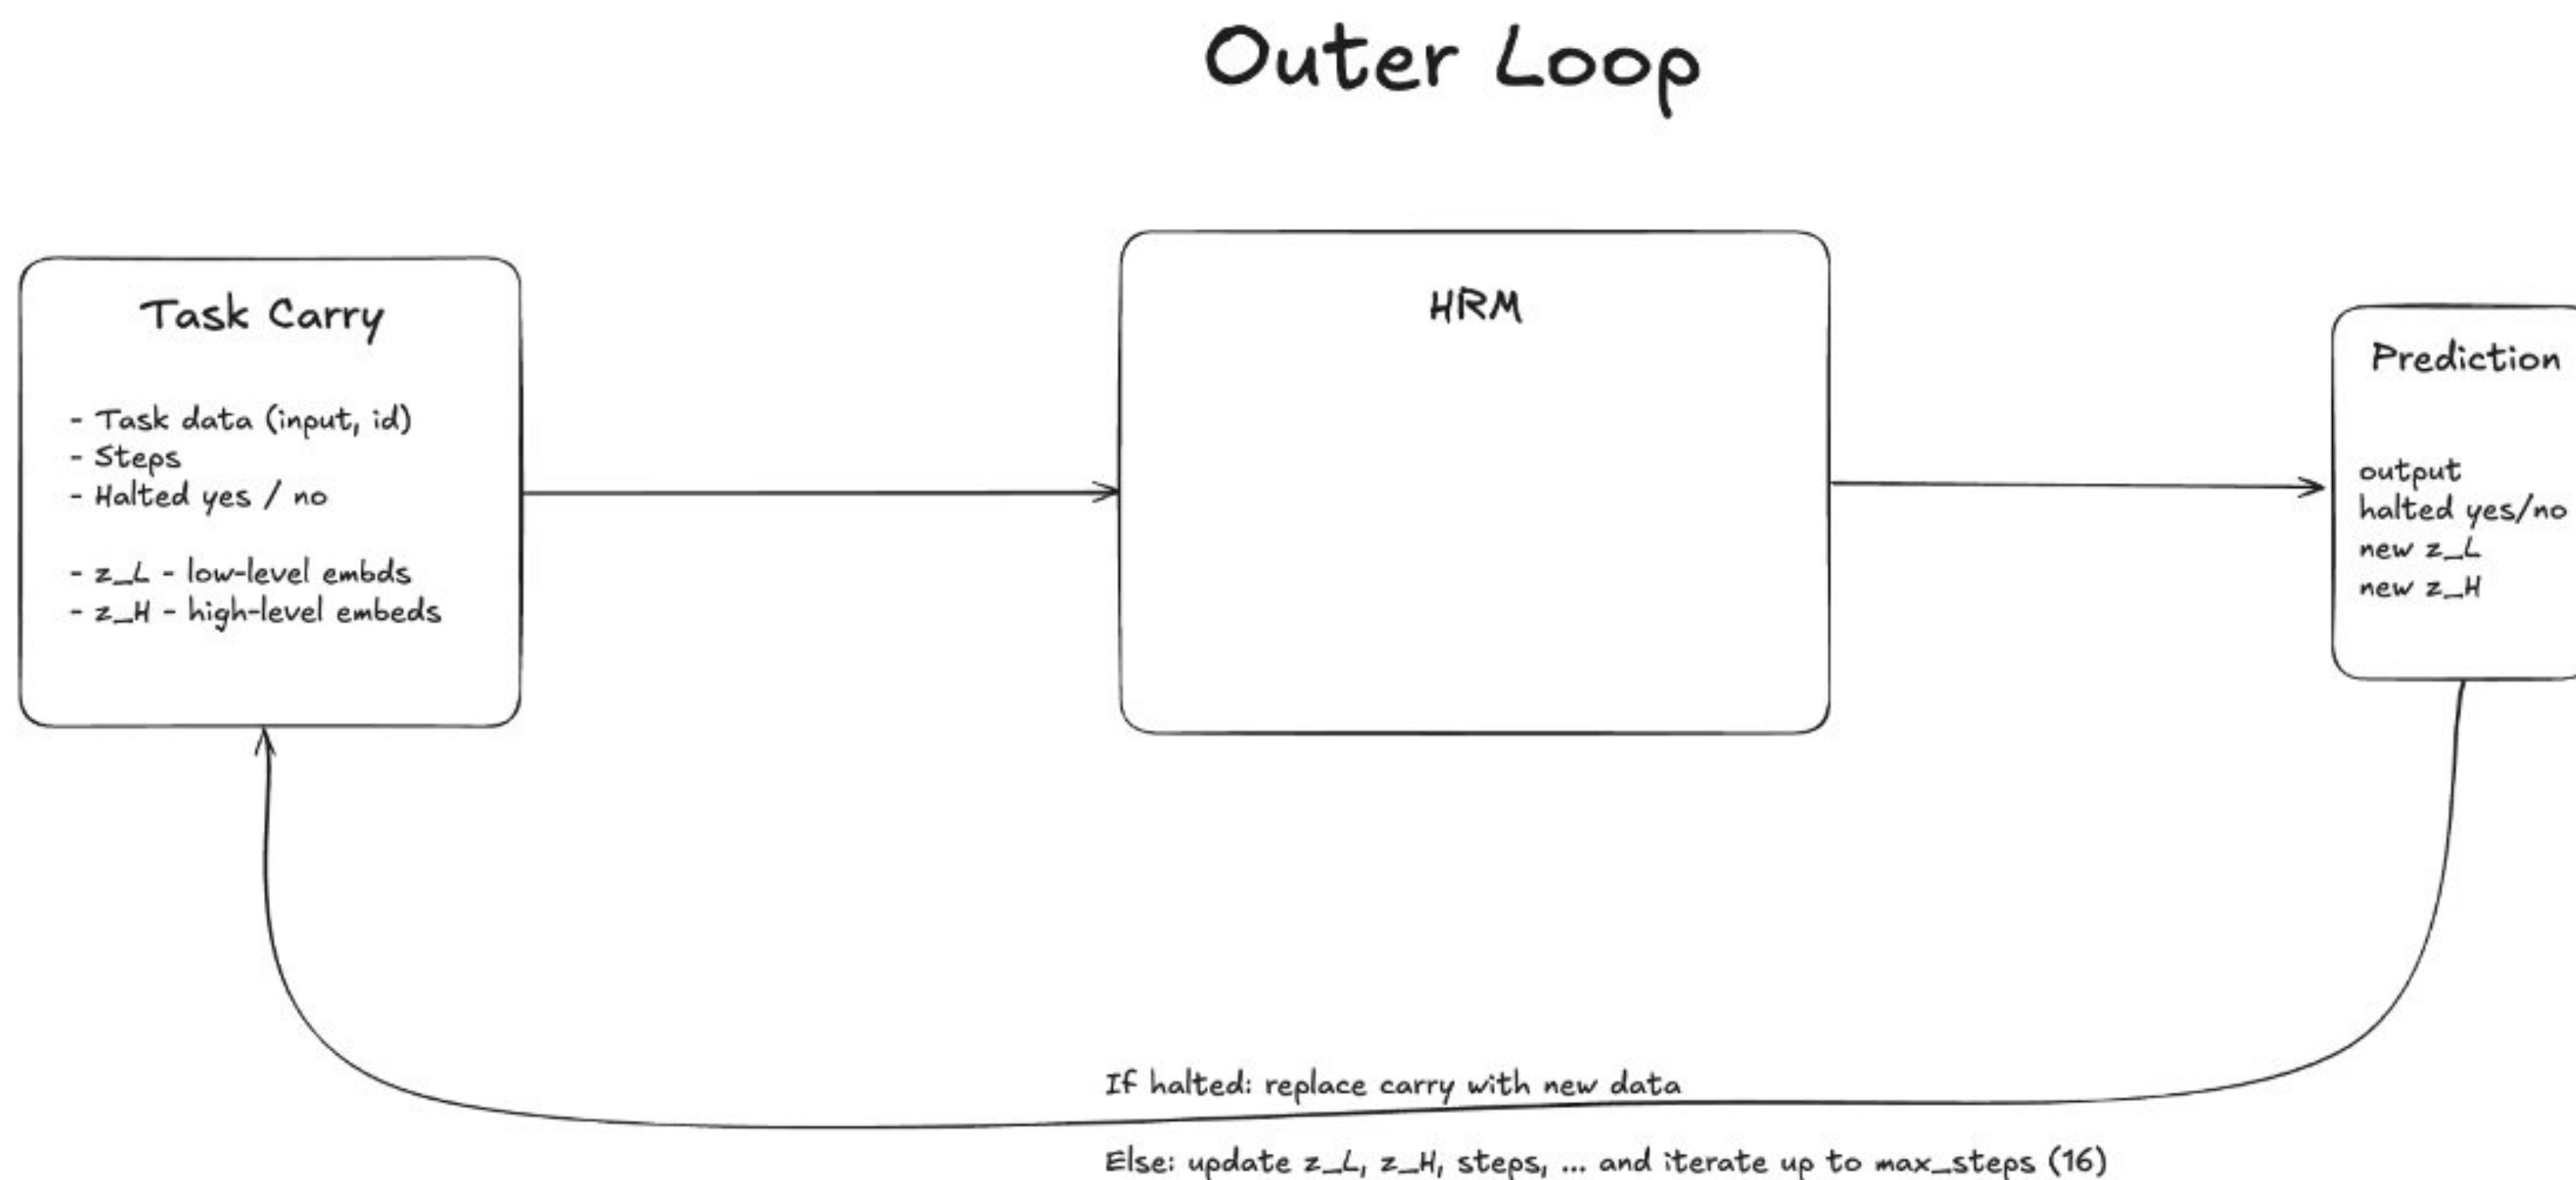
# Hierarchical Reasoning Model

## Overview

- Paper: <https://arxiv.org/abs/2506.21734>
- Code: <https://github.com/sapientinc/HRM>
- Analysis blog: <https://arcprize.org/blog/hrm-analysis>
- Analysis code: <https://github.com/arcprize/hierarchical-reasoning-model-analysis>

# Model Architecture

## Hierarchical Reasoning Model - Methods

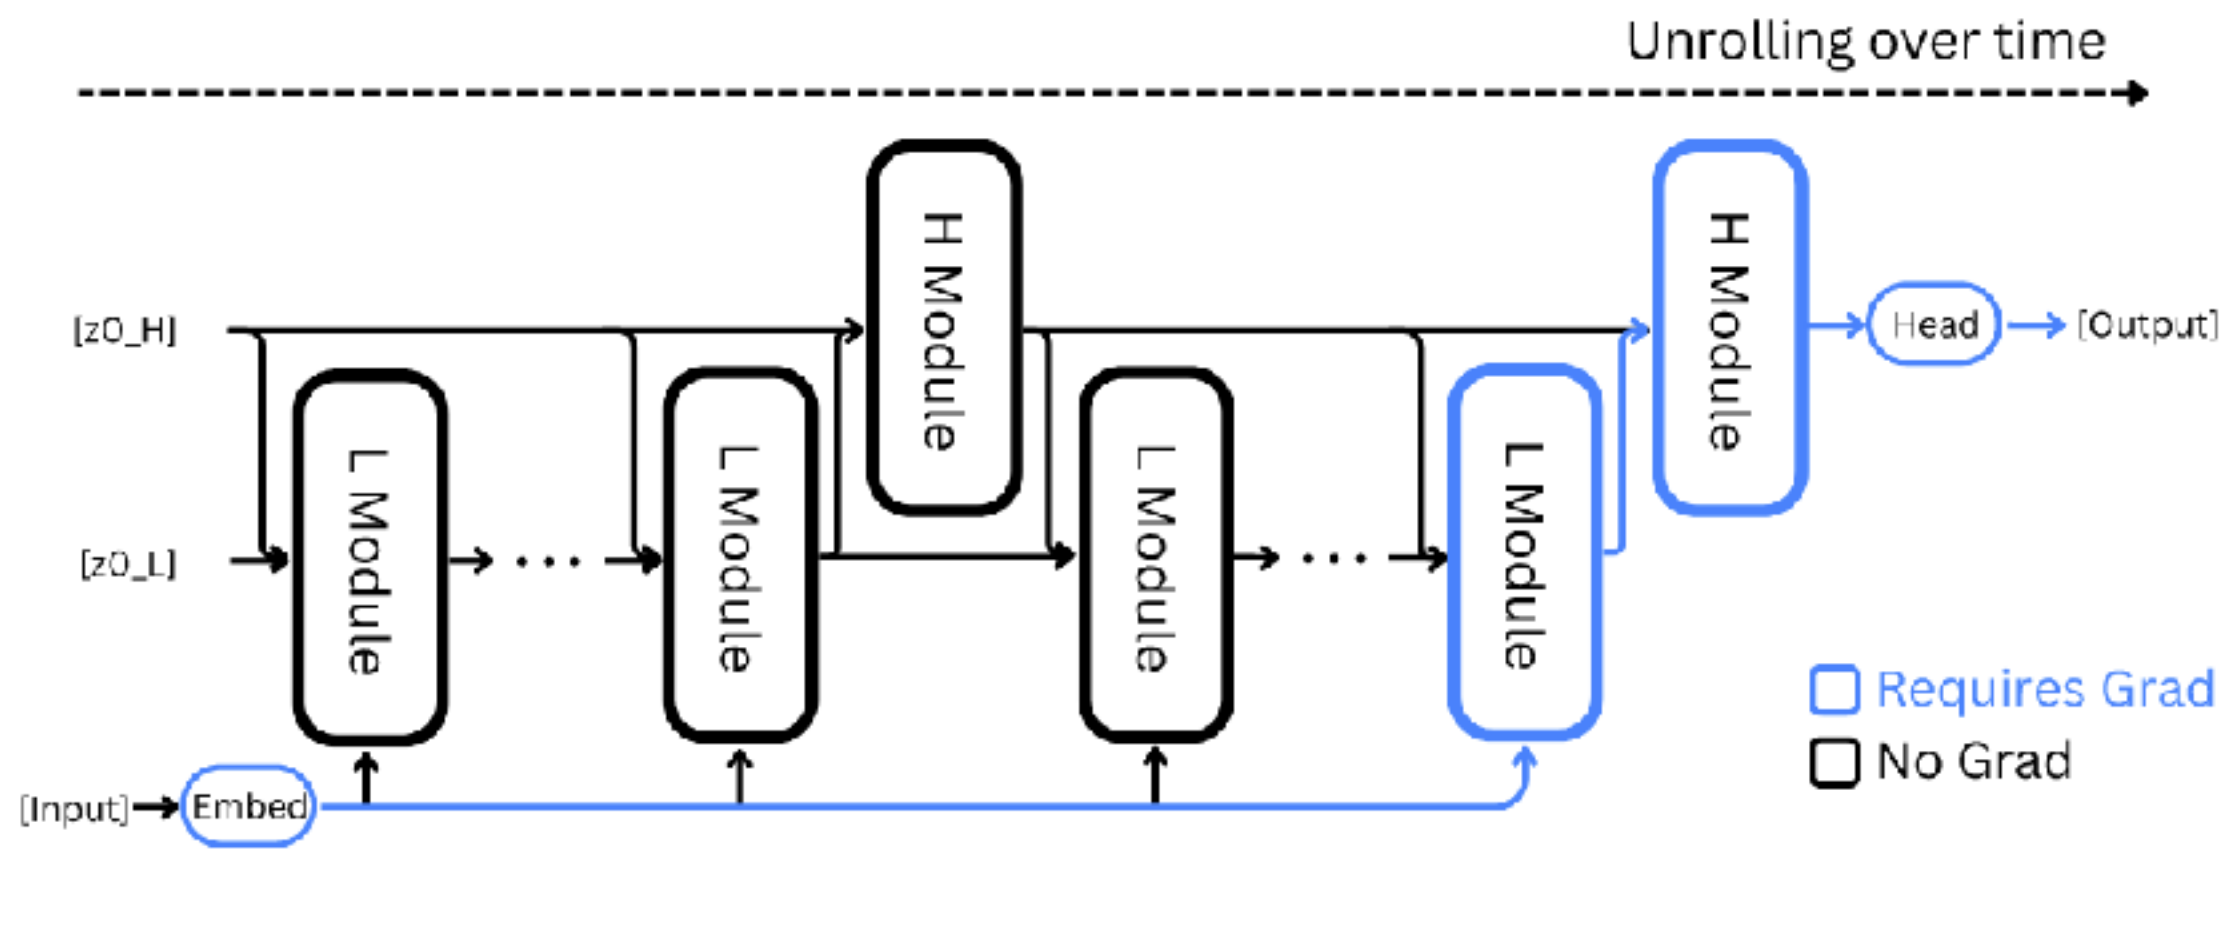


### Iterative Refinement

- given a context + initial guess, refine towards equilibrium
- Outer loop + 2 levels of inner loop
- Deep supervision (segment)

# Model Architecture

## Hierarchical Reasoning Model - Methods



$$\tilde{x} = f_I(x; \theta_I) \quad \hat{y} = f_O(z_H^{NT}; \theta_O)$$

```
def hrm(z, x, N=2, T=2):
    x = input_embedding(x)
    zH, zL = z

    with torch.no_grad():
        for _i in range(N * T - 1):
            zL = L_net(zL, zH, x)
            if (_i + 1) % T == 0:
                zH = H_net(zH, zL)

    # 1-step grad
    zL = L_net(zL, zH, x)
    zH = H_net(zH, zL)
    return (zH, zL), output_head(zH)

# Deep Supervision
for x, y_true in train_dataloader:
    z = z_init
    for step in range(N_supervision):
        z, y_hat = hrm(z, x)

        loss = softmax_cross_entropy(y_hat, y_true)
        z = z.detach()

    loss.backward()
    opt.step()
    opt.zero_grad()
```

**Low-level:**  $T$  time-steps each cycle

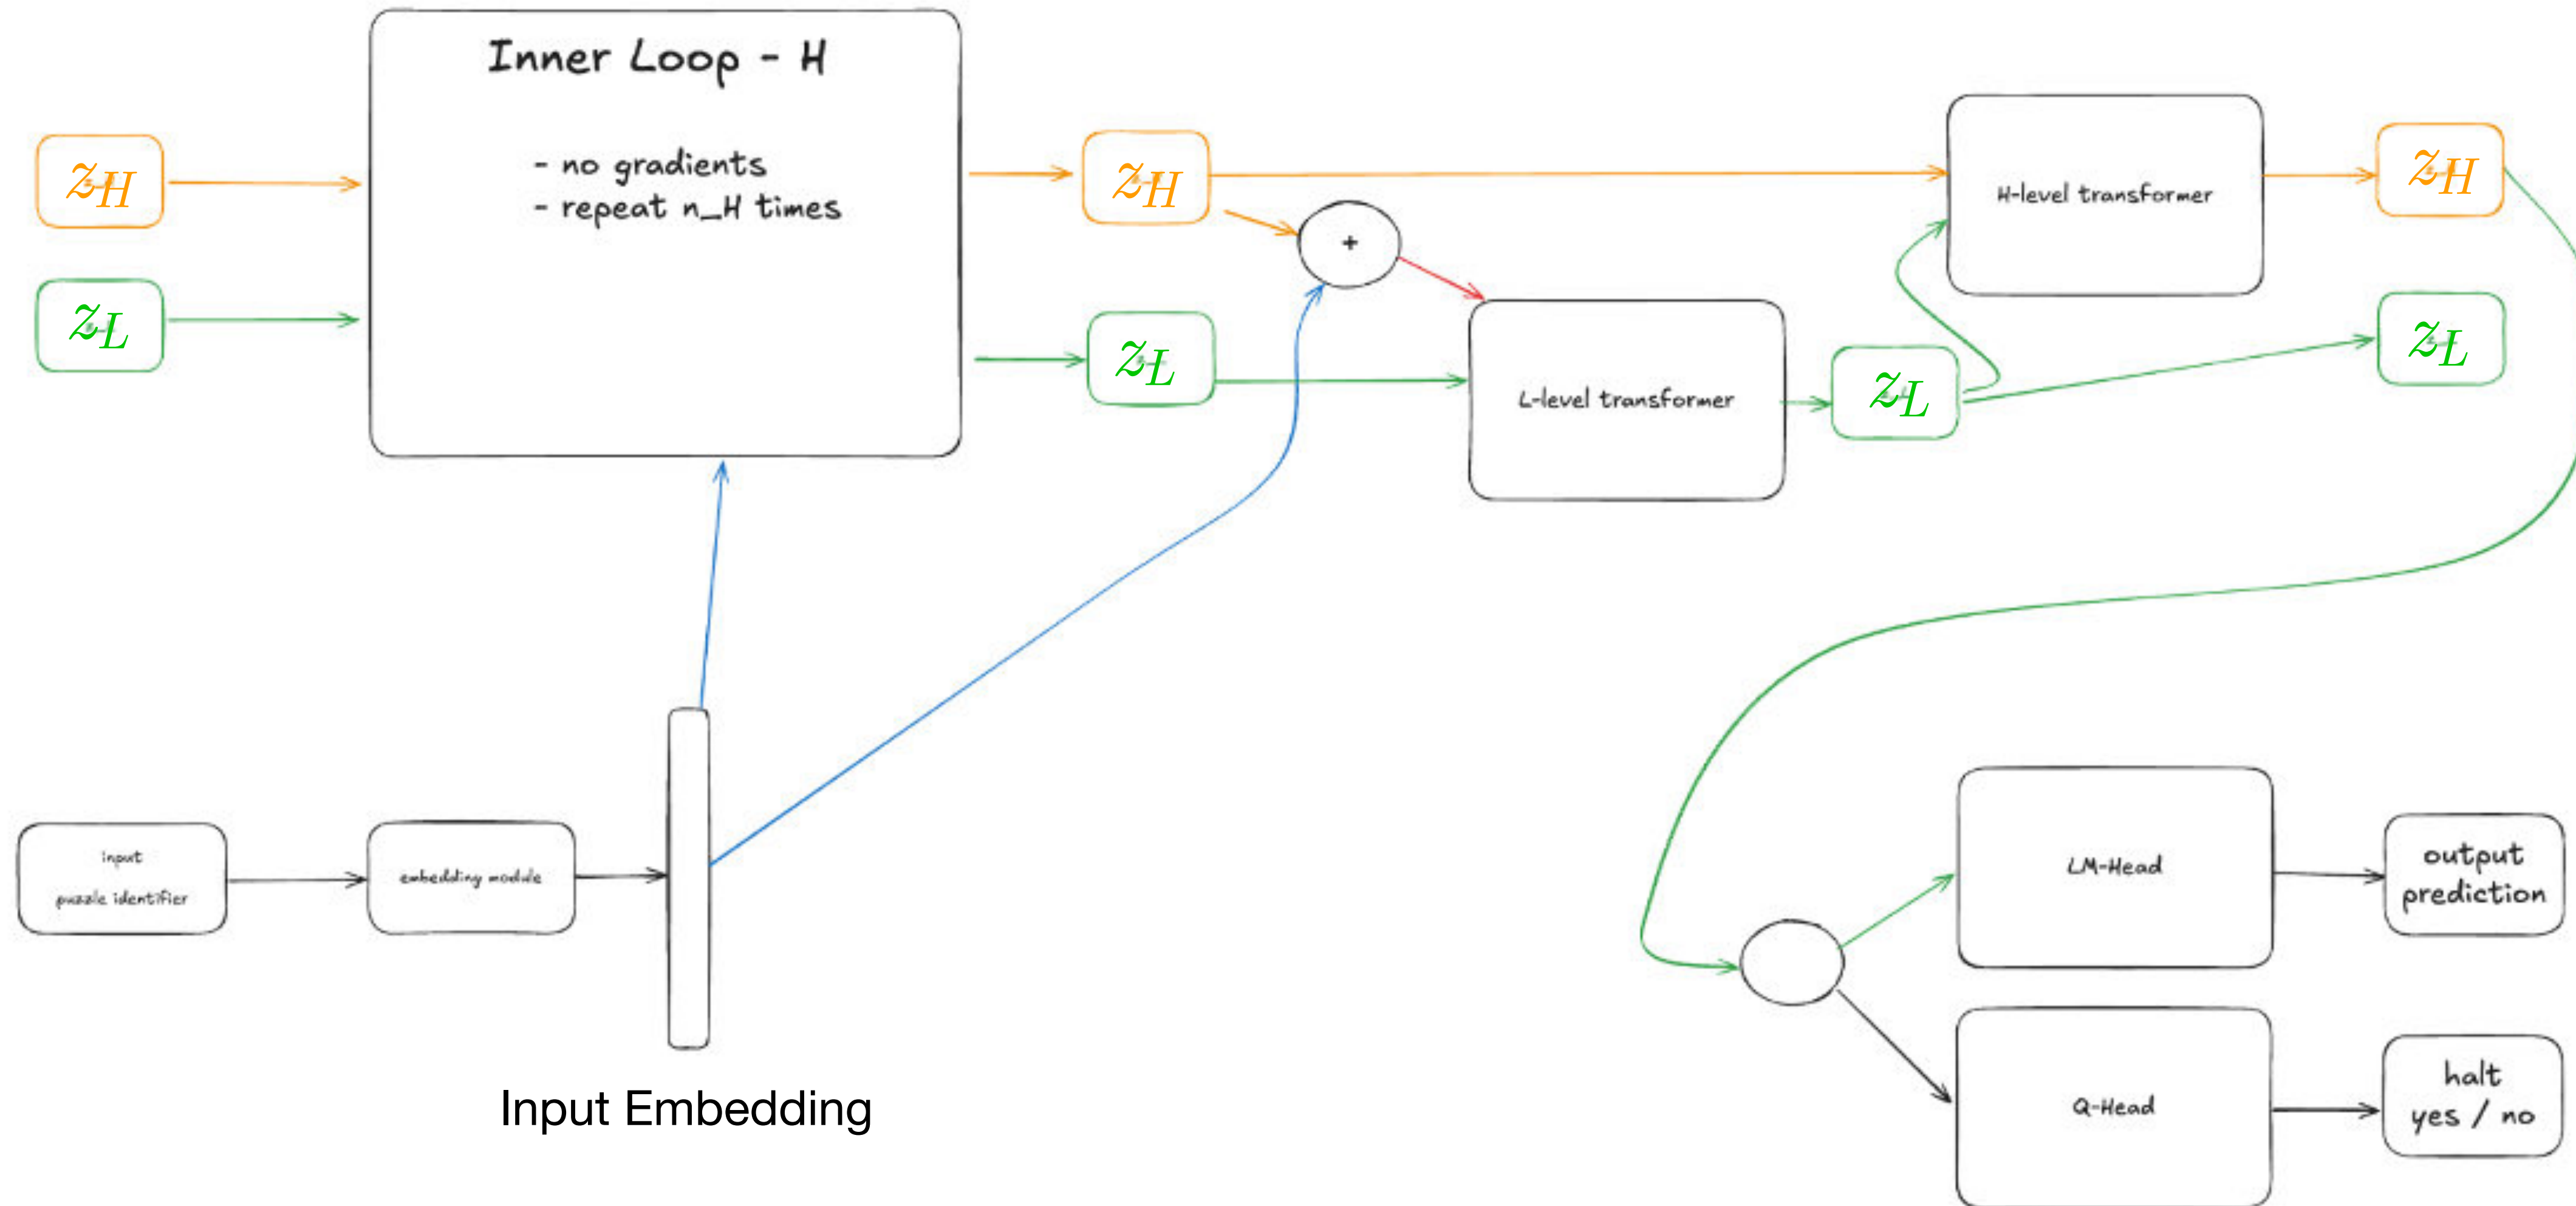
$$z_L^i = f_L(z_L^{i-1}, z_H^{i-1}, \tilde{x}; \theta_L)$$

**High-level:** executed once per cycle for  $N$  cycle


$$z_H^i = \begin{cases} f_H(z_L^{i-1}, z_H^{i-1}, \tilde{x}; \theta_H) & , \text{if } i \equiv 0 \pmod{T}; \\ z_H^{i-1} & , \text{otherwise.} \end{cases}$$



# HRM



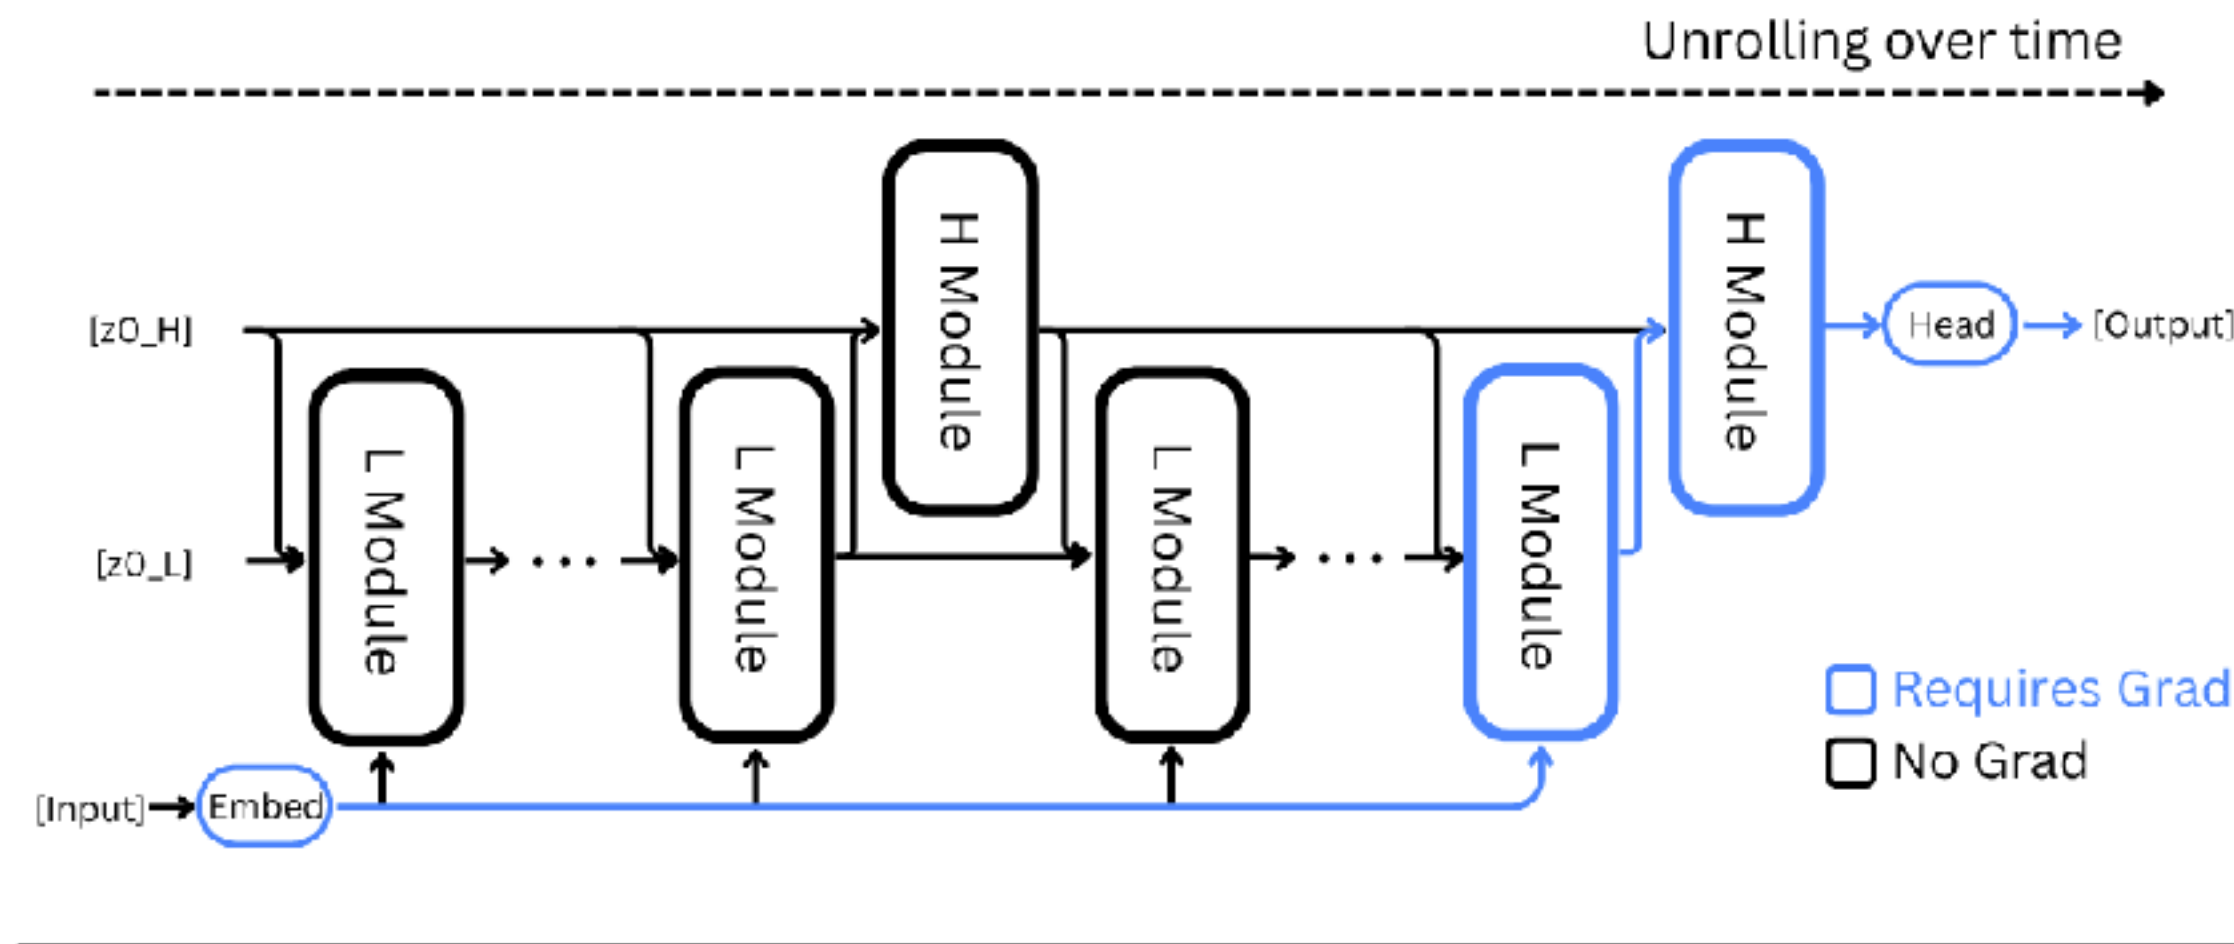
# Approximate gradient Hierarchical Reasoning Model



$$L_{\text{ACT}}^m = \text{LOSS}(\hat{y}^m, y) + \text{BINARYCROSSENTROPY}(\hat{Q}^m, \hat{G}^m) .$$

$$(I - J_f) \frac{dz^*}{d\theta} = \frac{\partial f}{\partial \theta}$$

$$(I - J_f)^{-1} = I + J_f + J_f^2 + \dots \approx I$$



One-Step gradient approximation

$$\frac{\partial z_H^*}{\partial \theta_H} \approx \frac{\partial f_H}{\partial \theta_H}, \quad \frac{\partial z_H^*}{\partial \theta_L} \approx \frac{\partial f_H}{\partial z_L^*} \cdot \frac{\partial z_L^*}{\partial \theta_L}, \quad \frac{\partial z_H^*}{\partial \theta_I} \approx \frac{\partial f_H}{\partial z_L^*} \cdot \frac{\partial z_L^*}{\partial \theta_I}, \quad \frac{\partial z_L^*}{\partial \theta_L} \approx \frac{\partial f_L}{\partial \theta_L}, \quad \frac{\partial z_L^*}{\partial \theta_I} \approx \frac{\partial f_L}{\partial \theta_I} .$$

# Adaptive computational time (ACT)

## Hierarchical Reasoning Model - Methods

$$L_{\text{ACT}}^m = \text{LOSS}(\hat{y}^m, y) + \text{BINARYCROSSENTROPY}(\hat{Q}^m, \hat{G}^m) .$$

- **Q-learning: Markov Decision Process (MDP)**
- action = {halt, continue}, state =  $\{z_H^m\}$
- Reward: halt  $\rightarrow$  1 if predicts correct; else 0.      continue  $\rightarrow$  0.
- Q-value is predicted by a Q-head:  $\hat{Q}^m = \sigma(\theta_Q^\top z_H^{mNT}) = (\hat{Q}_{\text{halt}}^m, \hat{Q}_{\text{continue}}^m)$
- According to Bellman Equation, the target Q-value function:  $\hat{G}^m = (\hat{G}_{\text{halt}}^m, \hat{G}_{\text{continue}}^m)$

$$\hat{G}_{\text{halt}}^m = \mathbf{1}\{\hat{y}^m = y\}, \quad \hat{G}_{\text{continue}}^m = \begin{cases} \hat{Q}_{\text{halt}}^{m+1}, & \text{if } m \geq N_{\text{max}}, \\ \max(\hat{Q}_{\text{halt}}^{m+1}, \hat{Q}_{\text{continue}}^{m+1}), & \text{otherwise.} \end{cases}$$



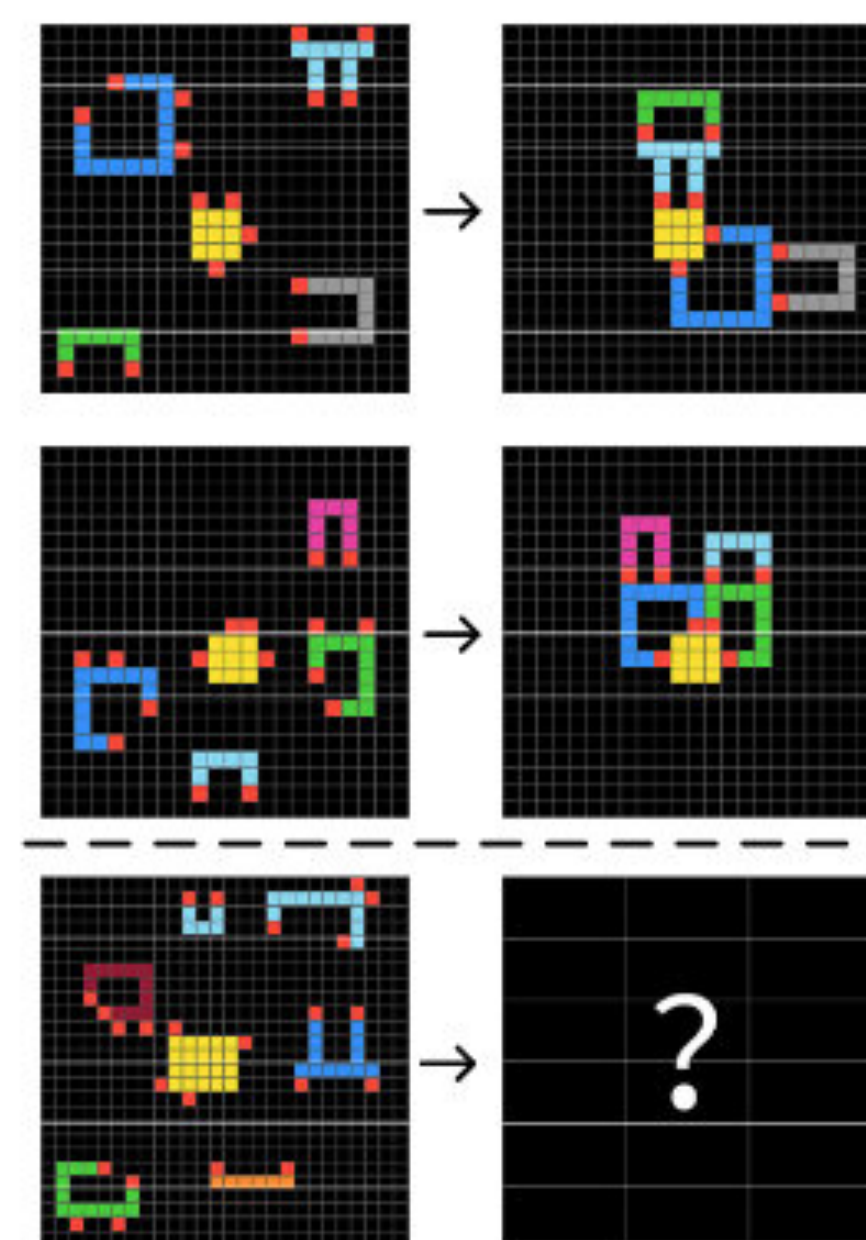
# Method Summary

## Hierarchical Reasoning Model

- **Outer Refinement Loop:** Propose solution and understand when you're done  
-> adaptive computation
- **Inner Equilibrium Loop:** 'hierarchical' refinement towards equilibrium
- **Puzzle Embeddings:** learn to associate a puzzle hash with a specific transformation

# Benchmarks

## Hierarchical Reasoning Model



(a) ARC-AGI

	8	4			5		6		
				8		7			
3				4					
	3	8	4				2		
9		6			3			8	
				5				6	
					2				1
	2	5		3			8		

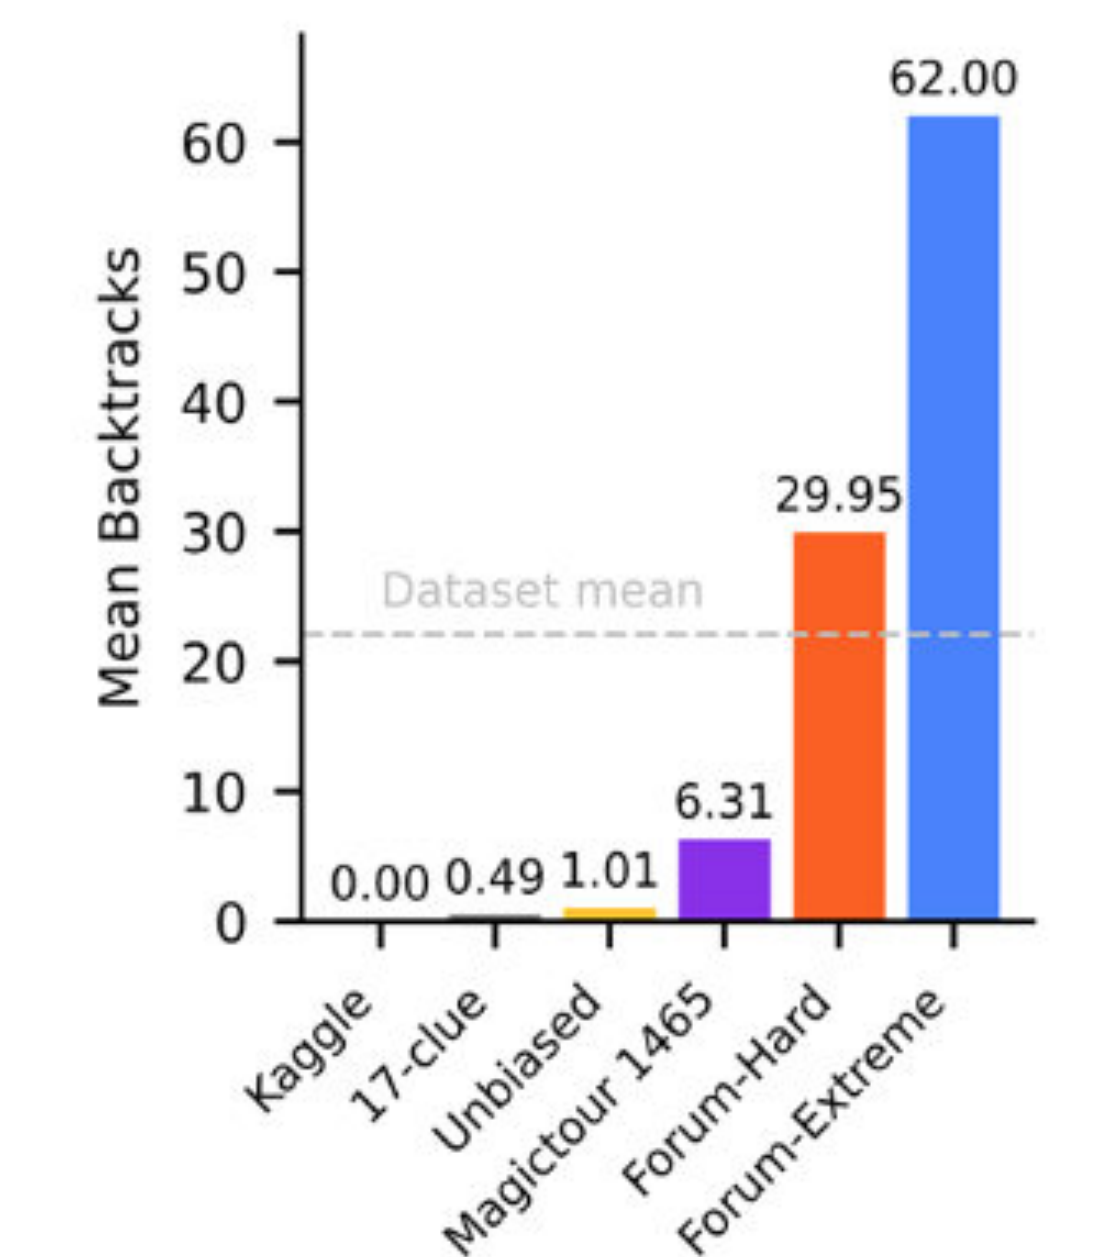
  

7	8	4	1	2	5	9	6	3	
2	6	1	3	8	9	7	4	5	
3	5	9	6	4	7	8	1	2	
5	3	8	4	9	6	1	2	7	
4	1	6	2	7	3	5	9	8	
9	7	2	8	5	1	4	3	6	
6	9	3	5	1	8	2	7	4	
8	4	7	9	6	2	3	5	1	
1	2	5	7	3	4	6	8	9	

(b) Sudoku-Hard



(c) Maze navigation

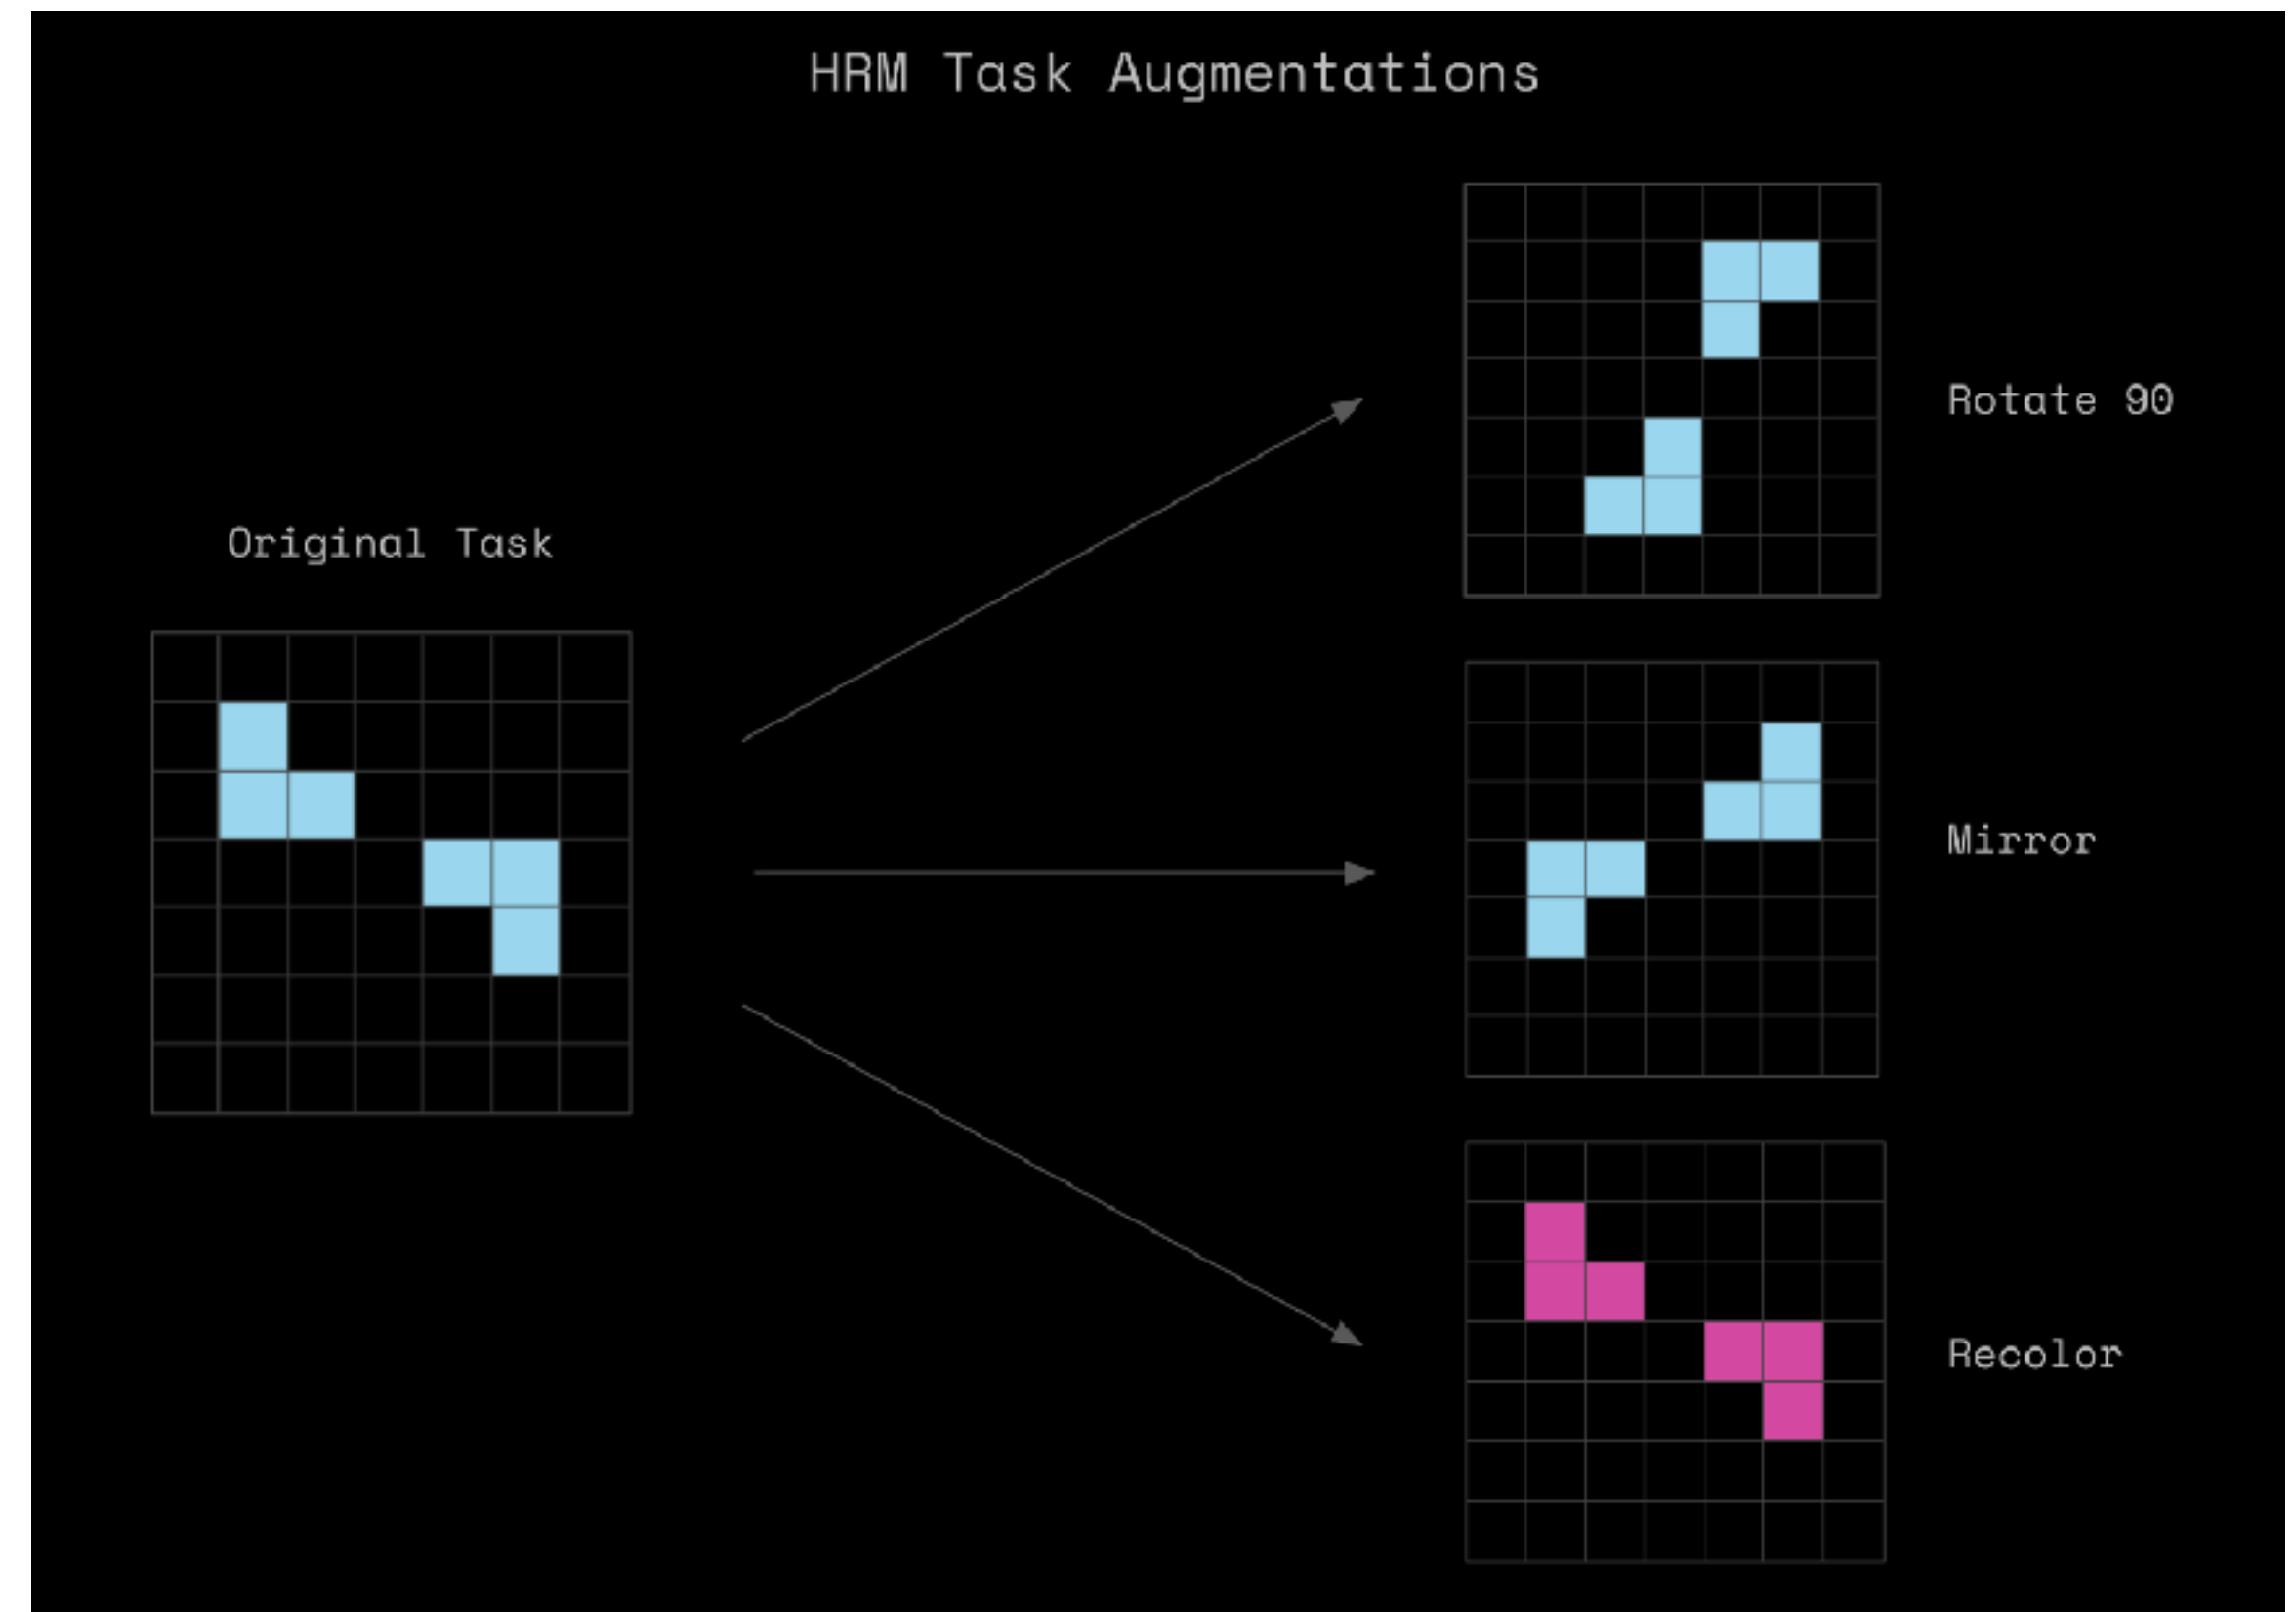


(d) *Sudoku-Extreme* subset difficulty

# Data Augmentation

## Hierarchical Reasoning Model

- Inference
  - Make separate predictions for all augmented versions of a task
  - Use maximum outer refinement loops, no ACT
  - Use majority voting on pool of predictions






# Performance Overview


## Hierarchical Reasoning Model

- Arc v1
  - Public 41 % pass@2
  - Semi-Private 32% pass@2
- Arc v2
  - Public 4% pass@2
  - Semi-Private 2% pass@2

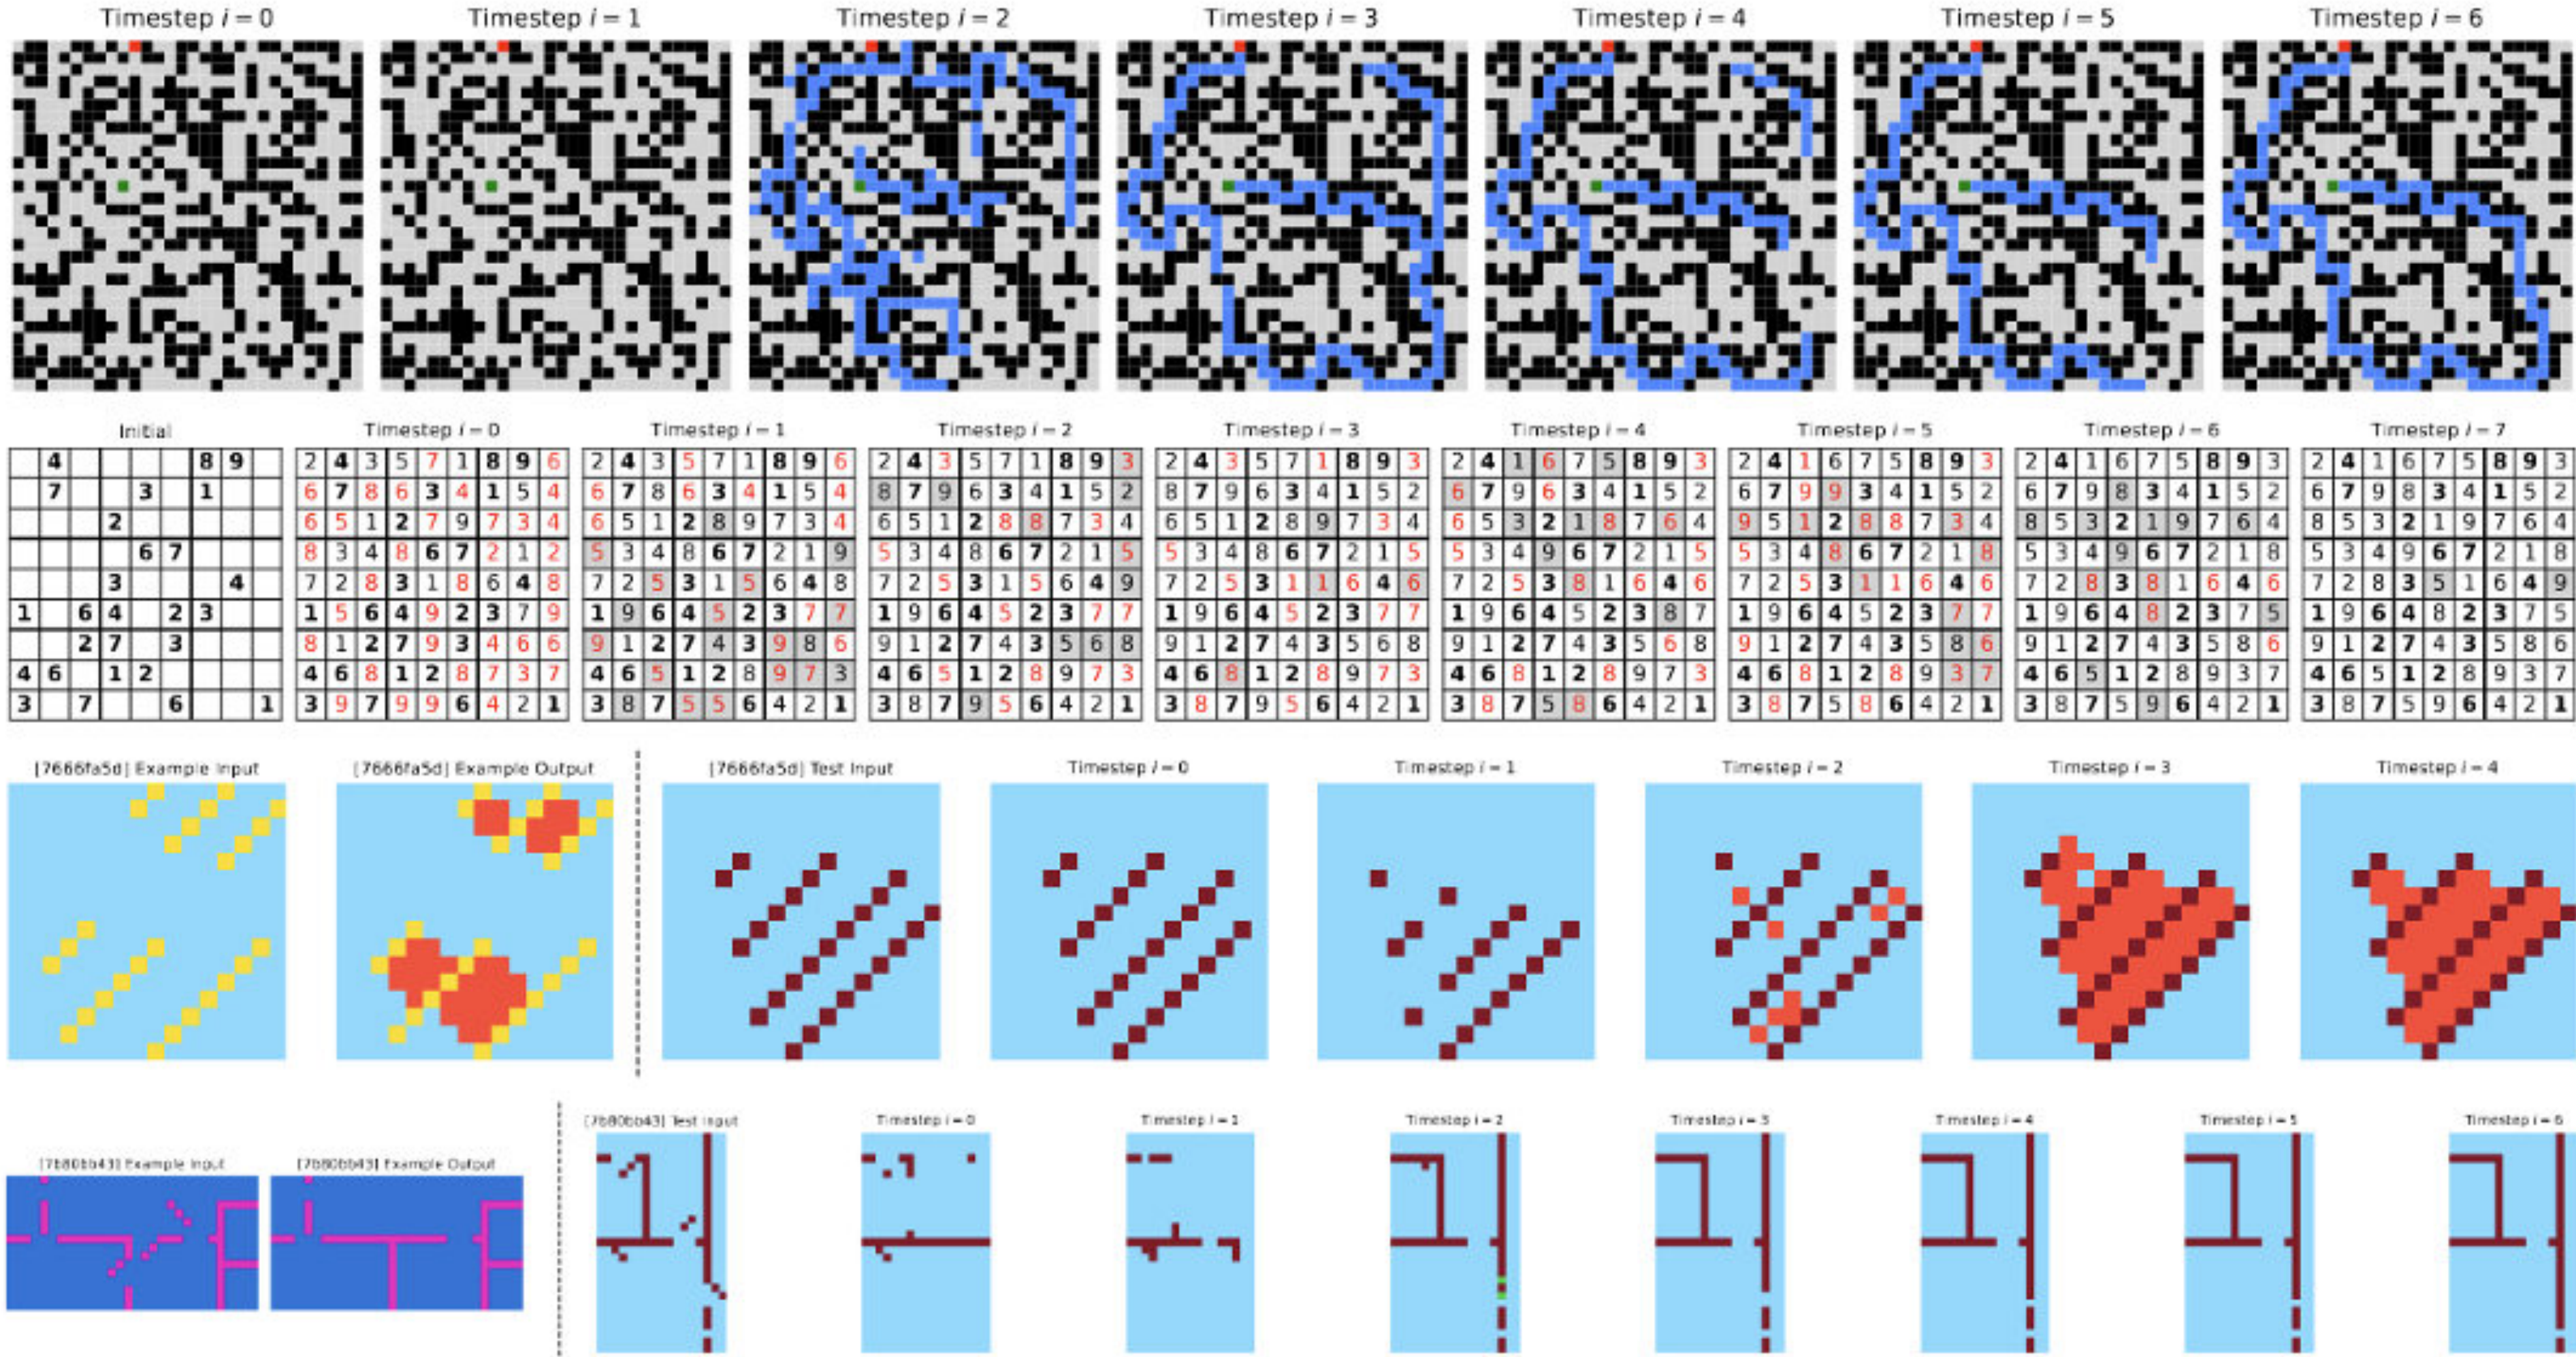


# Performance Overview

## Hierarchical Reasoning Model









# Why does it work?

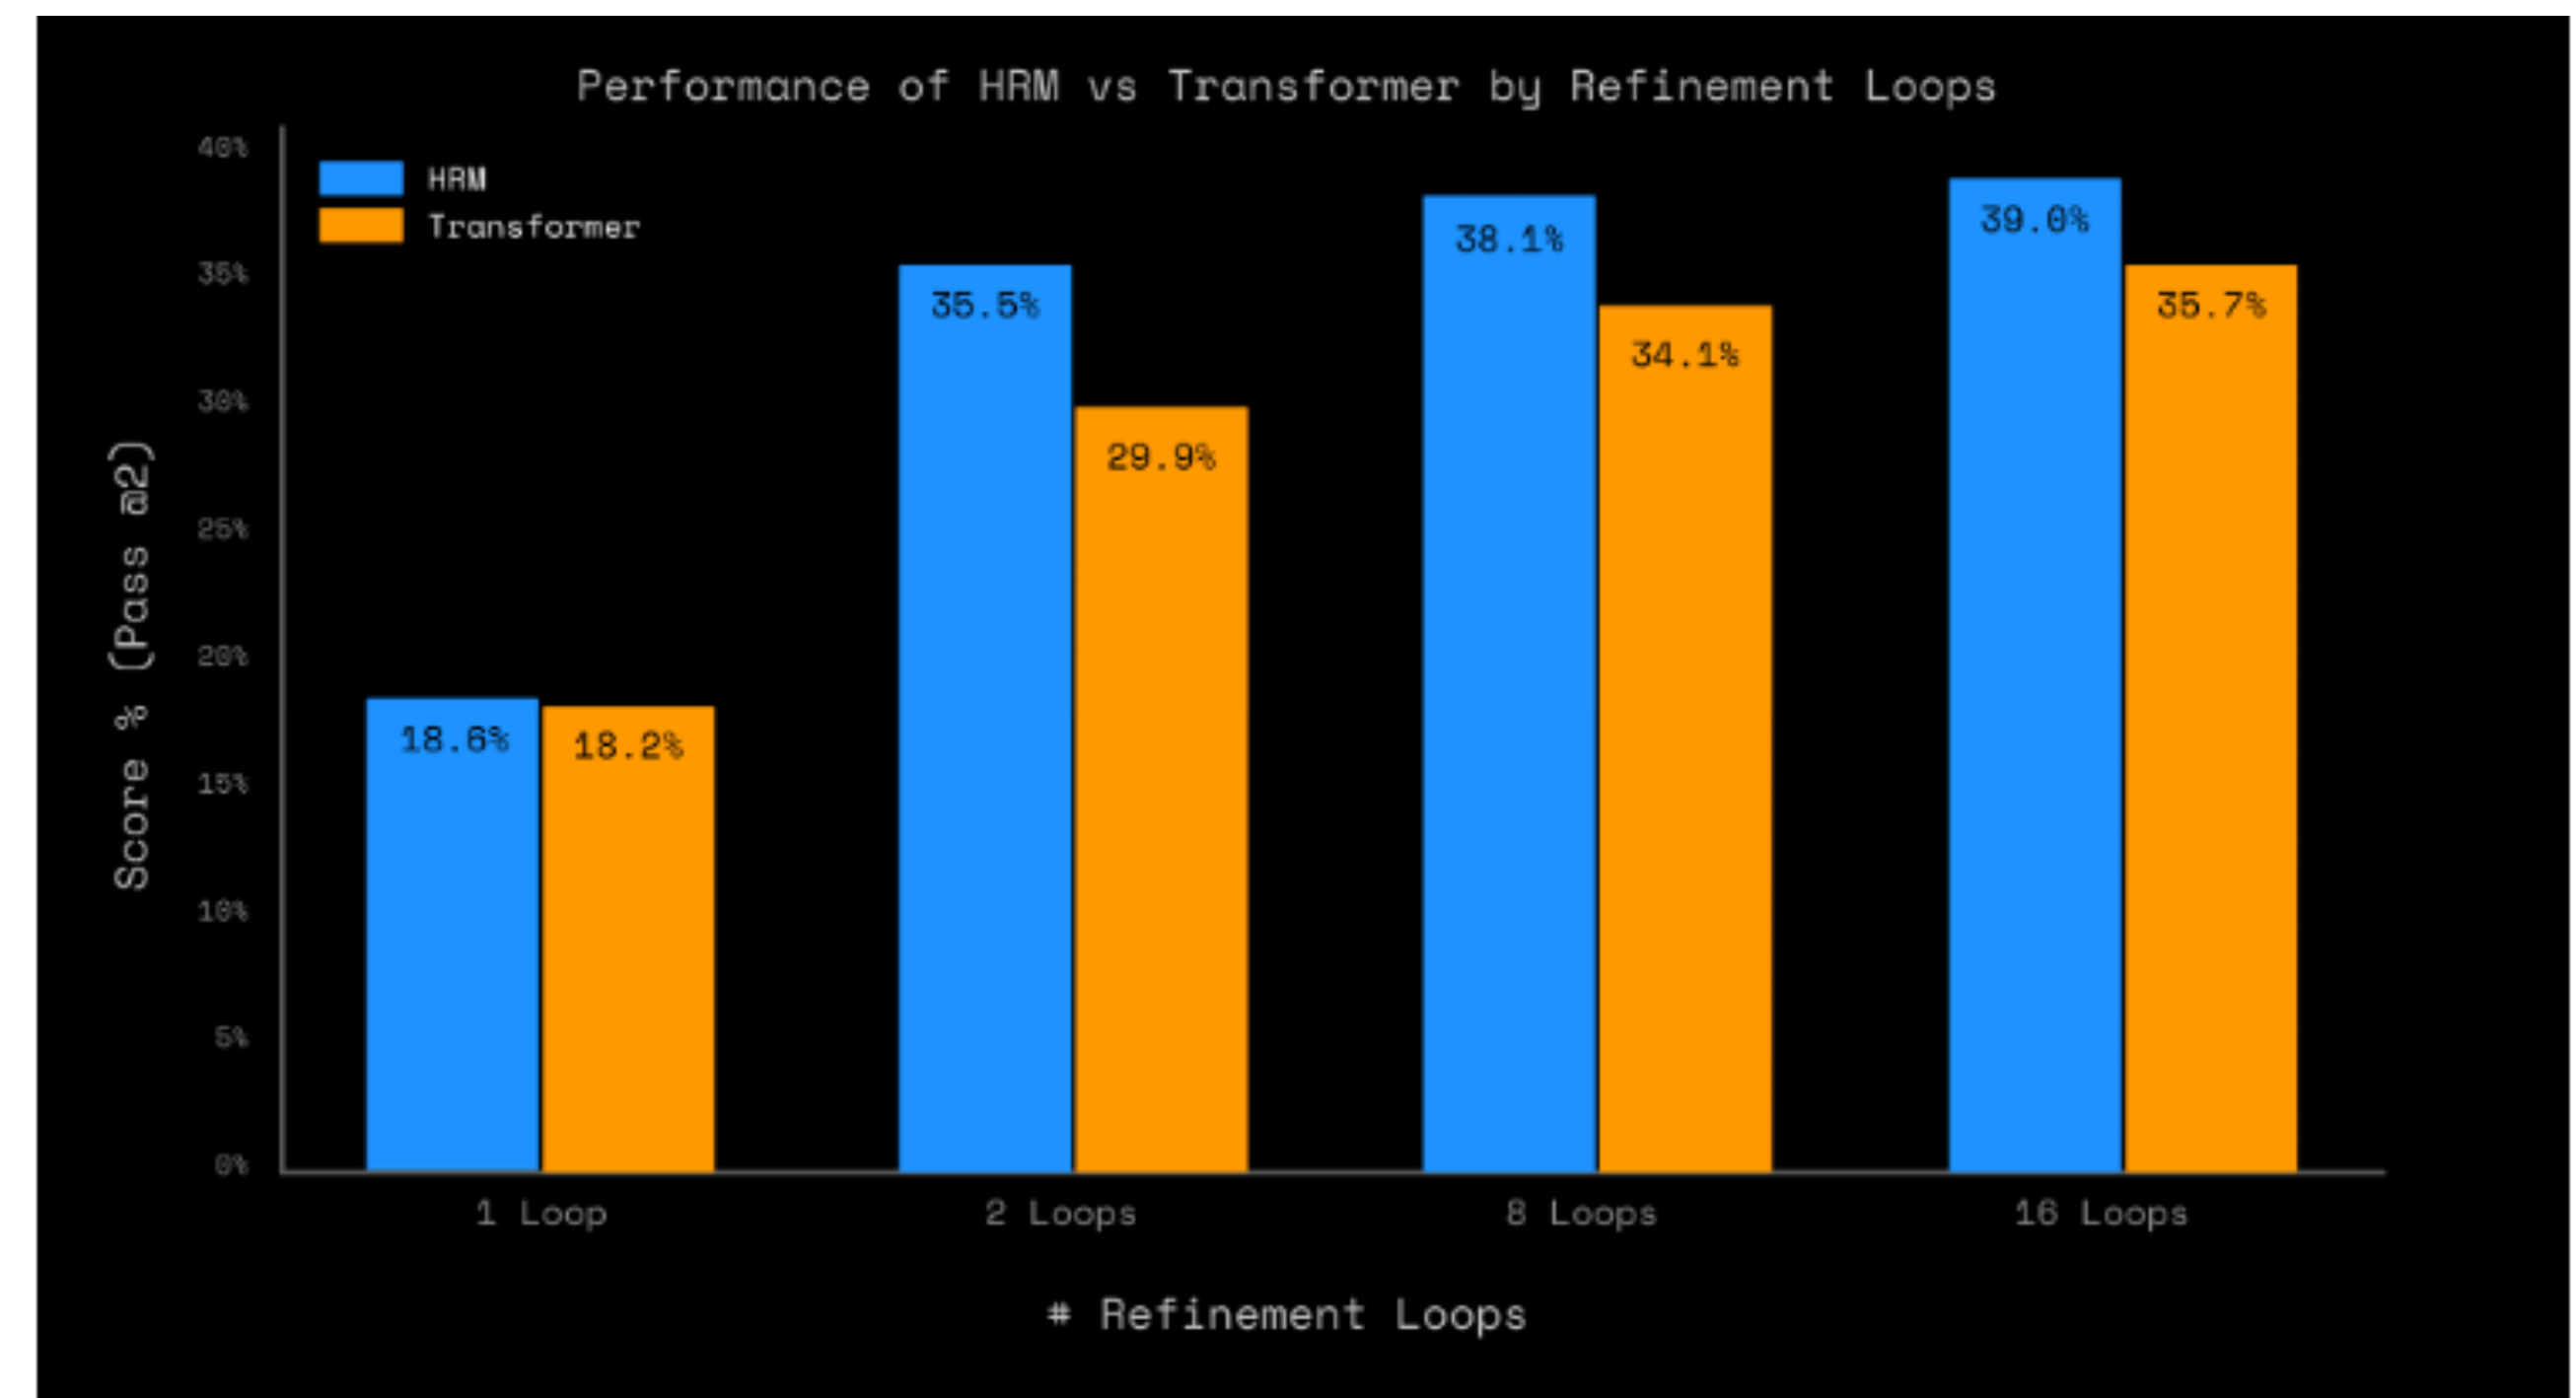
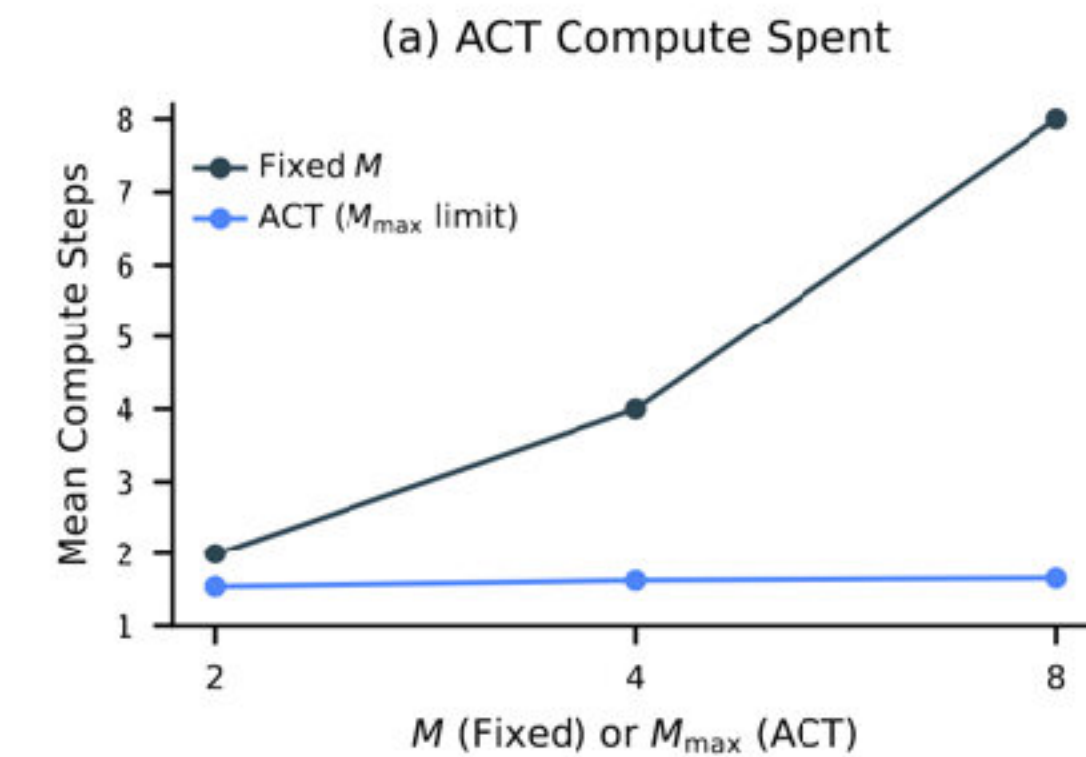
## Hierarchical Reasoning Model - Performance Drivers

- 1. Model architecture
- 2. Inner refinement loop
- 3. Outer refinement loop, with or without ACT
- 4. Data augmentation for HRM
- 5. Puzzle embeddings

# Architecture?

## Hierarchical Reasoning Model - Performance Drivers

- Compare HRM with regular transformer
  - Matched parameter count
  - Same hyperparameters
  - Less (forward) compute
- Findings:
  - Regular, unoptimized transformer is within ~5%.
  - No refinement -> matched
  - More outer loop steps -> gap closes







# Inner Loop?

## Hierarchical Reasoning Model - Performance Drivers

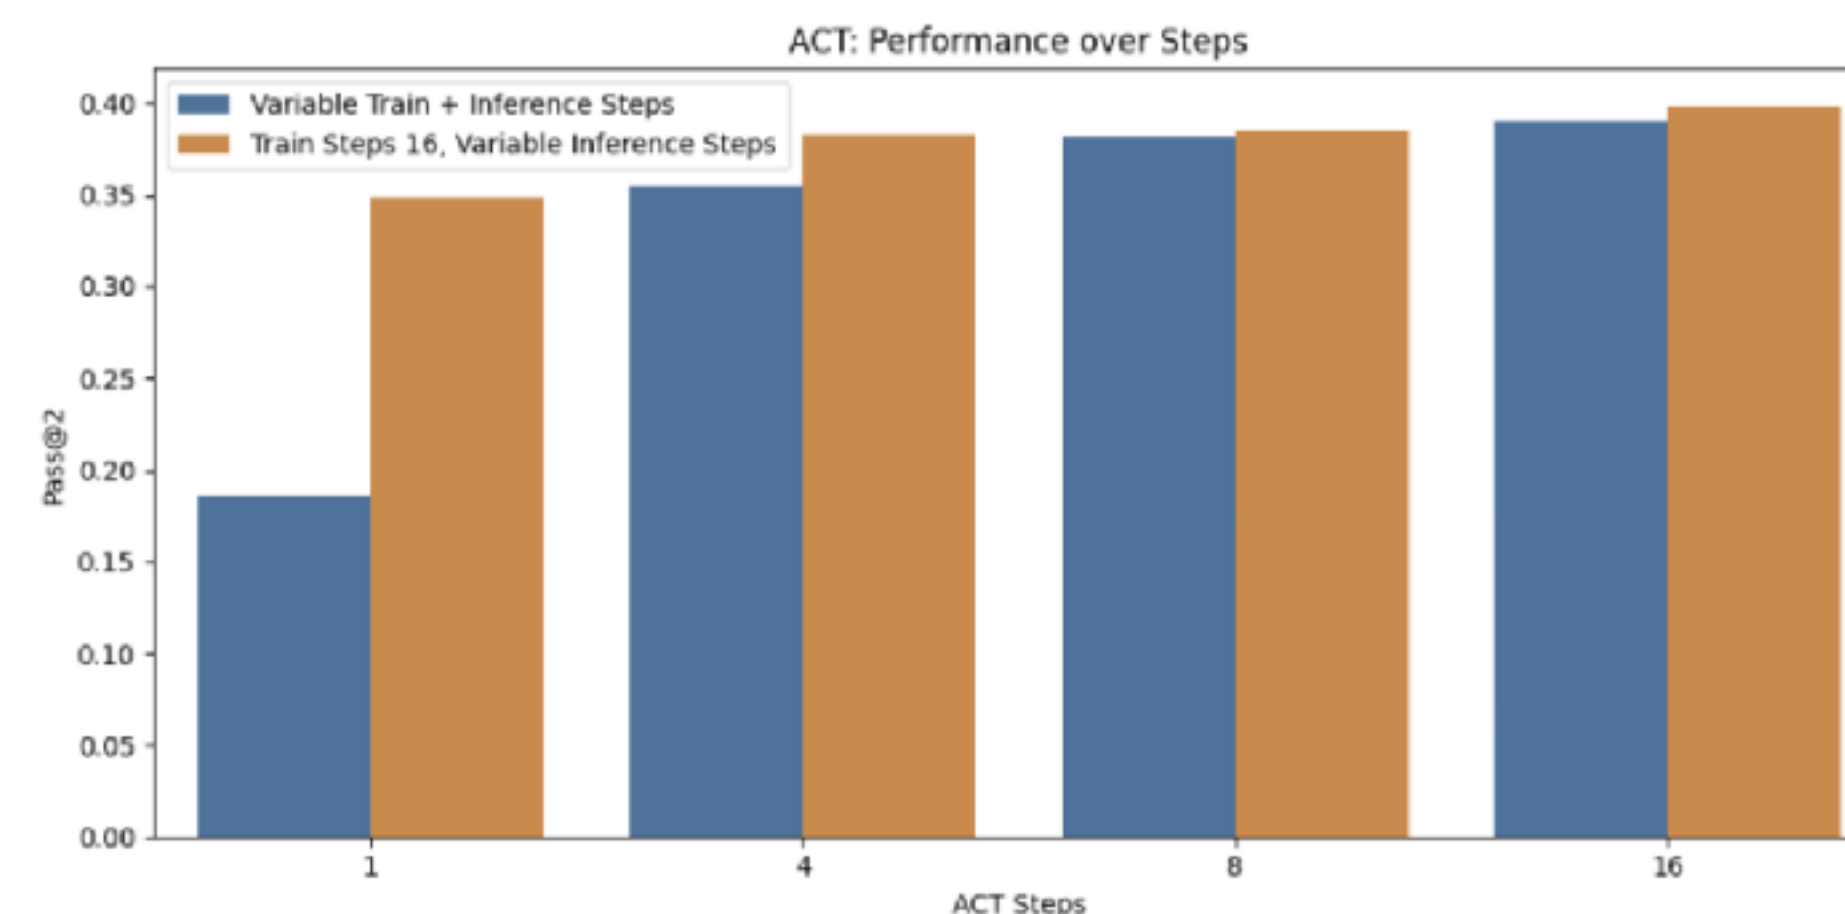
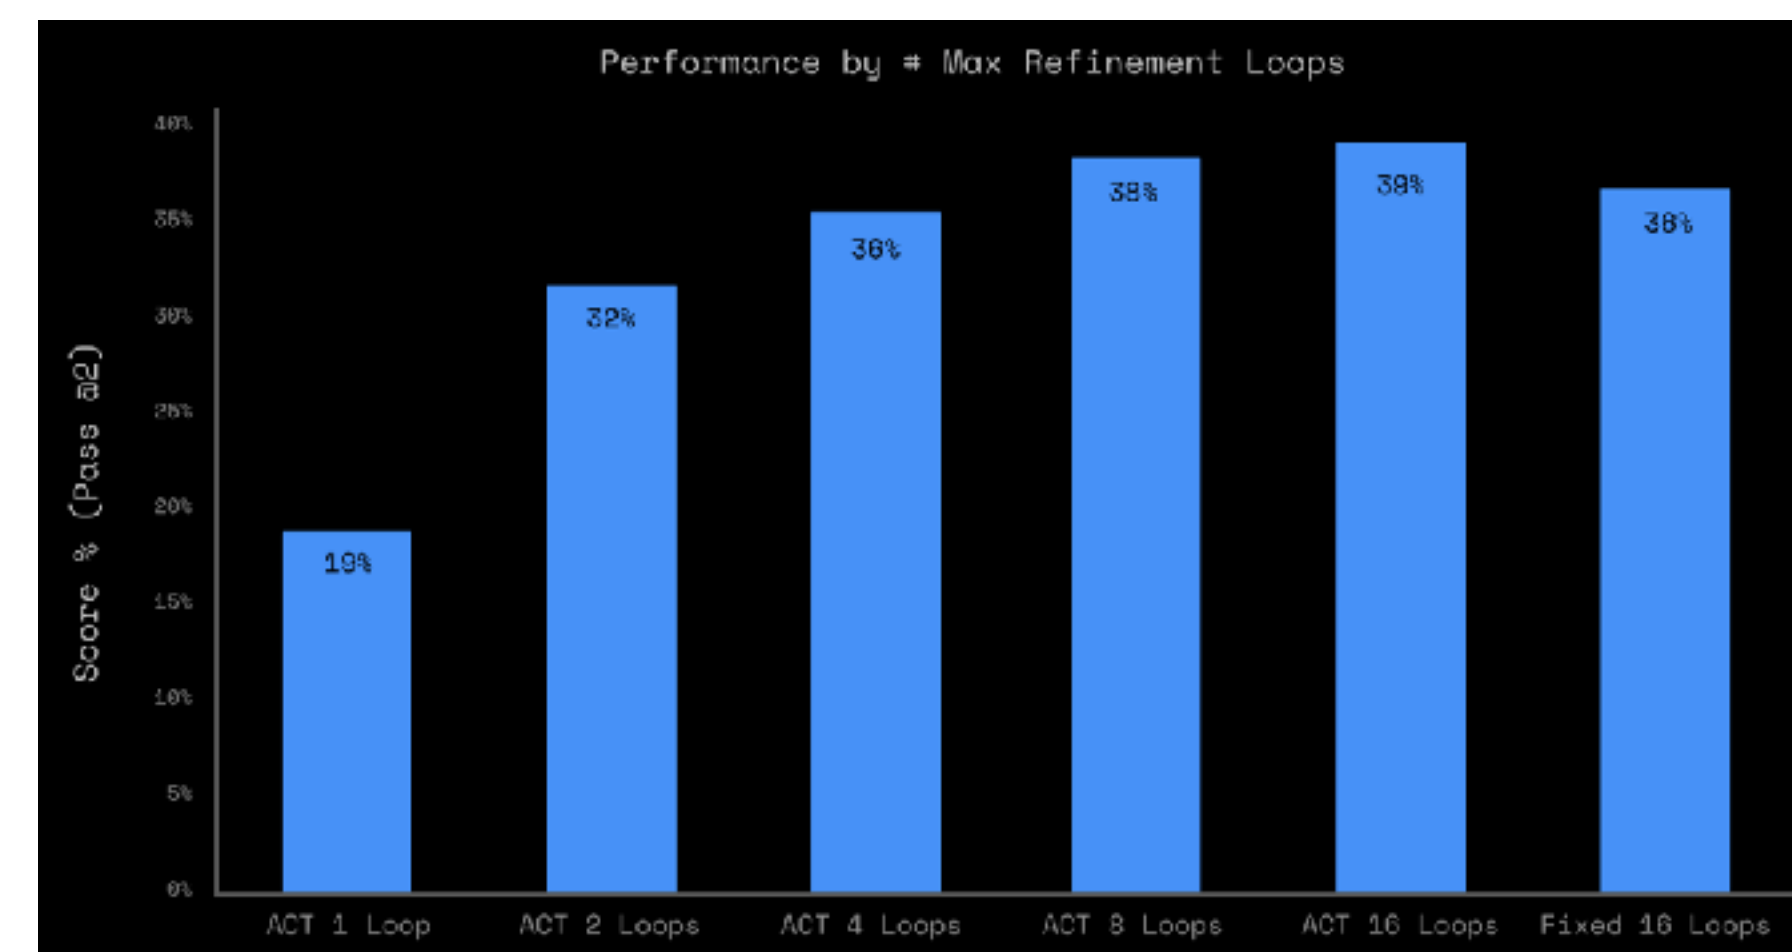
- Vary L and H cycles
- **Findings**
  - L1-H1 already achieves  $\frac{7}{8}$  of the performance.
  - Scaling higher does not improve further
  - #params remains the same, but compute scales with  $n_L * n_H$ .



# Outer Loop + ACT?

## Hierarchical Reasoning Model - Performance Drivers

- Vary max outer loop steps
- **Findings**
  - Refinement makes a big difference.
  - Refinement at training > refinement at inference
  - Learned halting signal leads works (only at training)
- **Refinement of predictions works!**



# Tiny Recursive Model

## Overview

- Paper: <https://arxiv.org/abs/2510.04871>
- Code: <https://github.com/SamsungSAILMontreal/TinyRecursiveModels>

# Motivations

## Improvements in Hierarchical Reasoning Models

- Implicit Function Theorem with 1-step gradient approximation  
Although the residuals do generally reduce over time, a fixed point is unlikely to be reached when the theorem is actually applied.
- Twice the forward passes with Adaptive computational time (ACT)
- Why ‘Hierarchical’ interpretation? a bit forced...

# Model Design

## Tiny Recursive Model

- No hierarchy! No IFT! No gradient approximation!
- Forget the biological story.
- $z_H$  = the current embedded solution. → Decoded via output head.
- $z_L$  = latent reasoning state, not directly a solution.

- Example: Sudoku —

- $z_H$  matches the current solution;
- $z_L$  does not.

5	2	6	7	9	4	8	3	1
3	9	1	2	6	8	4	7	5
4	8	7	3	1	5	2	9	6
1	6	8	5	3	2	7	4	9
9	3	5	4	7	6	1	8	2
7	4	2	9	8	1	5	6	3
8	7	3	1	5	9	6	2	4
2	5	9	6	4	7	3	1	8
6	1	4	8	5	3	9	5	7

Tokenized  $z_H$  (denoted  $y$  in TRM)

5		5	4	9	4		6	3
4		3	1			4	6	5
4	8	4		3		6	6	4
9		6	5	3		5	4	
	3	5	4	3		5	4	4
6		3		3	3	5	8	8
3	3	3	6	5		6	6	4
7	5		6		3	3	6	6
4	3	4	8		3	6	6	4

Tokenized  $z_L$  (denoted  $z$  in TRM)

$$z_L \leftarrow f_L(z_L + z_H + x)$$

⋮

$$z_L \leftarrow f_L(z_L + z_H + x)$$

$$z_H \leftarrow f_H(z_L + z_H)$$

# Model Design

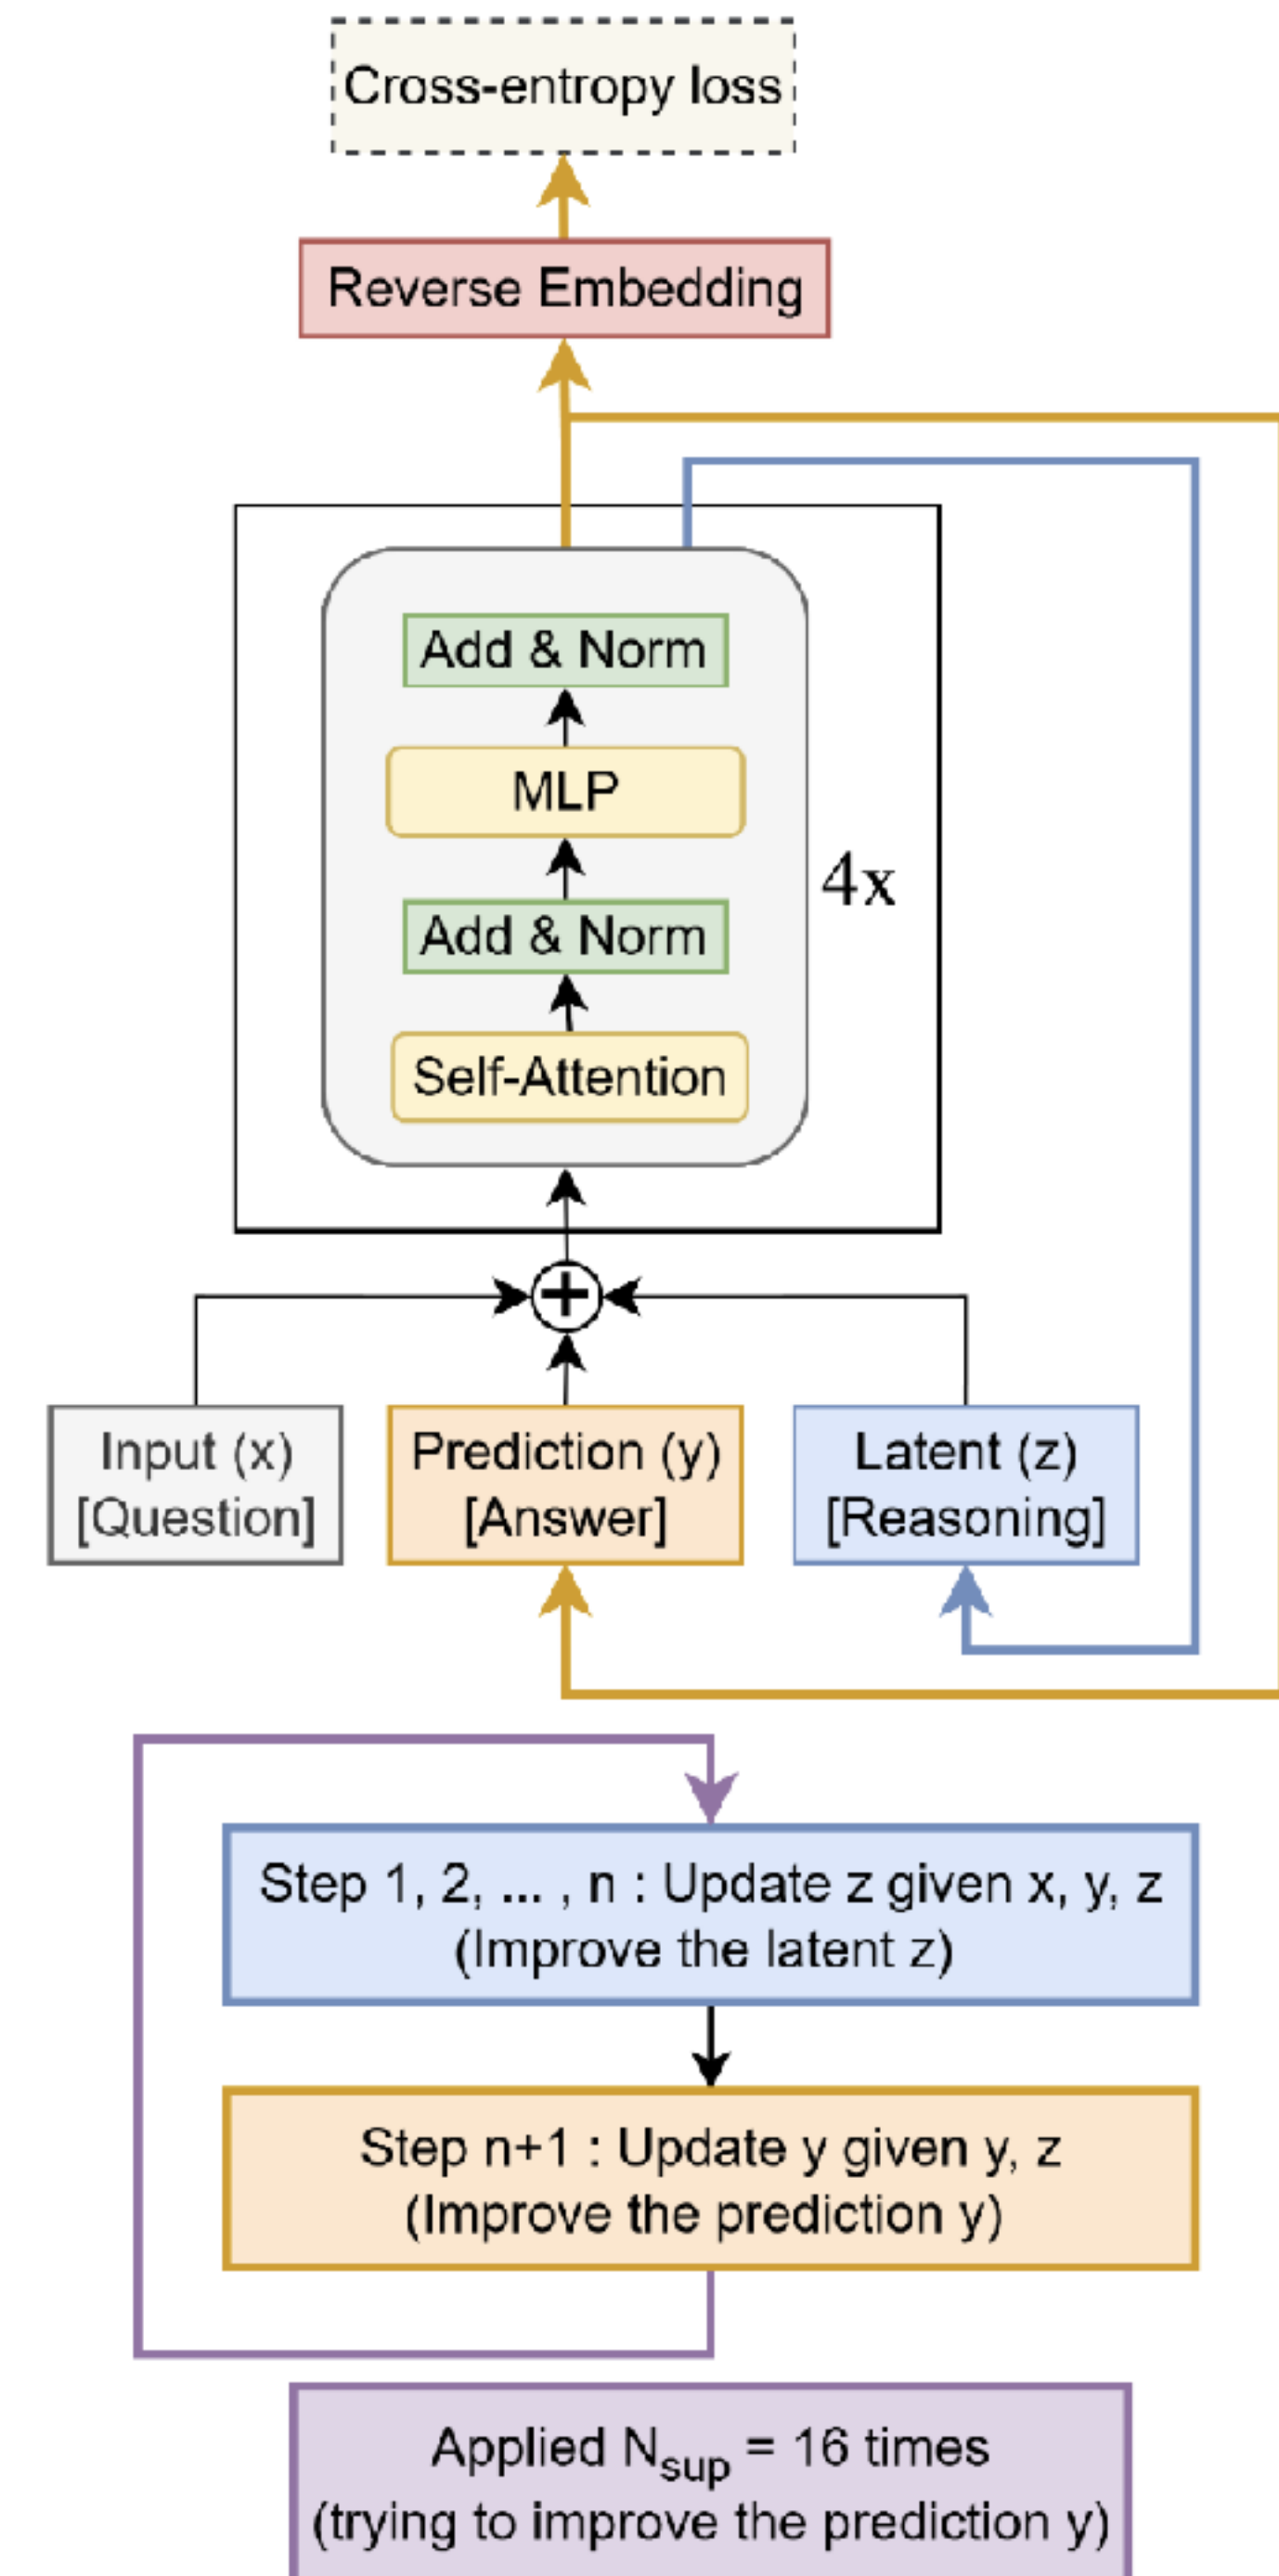
## Tiny Recursive Model

- Why 2 latent variables?  
a proposed solution  $y (z_H)$ , and a latent reasoning feature  $z (z_L)$ .

Table 2. TRM on Sudoku-Extreme comparing % Test accuracy when using more or less latent features

Method	# of features	Acc (%)
TRM $y, z$ (Ours)	2	87.4
TRM multi-scale $z$	$n + 1 = 7$	77.6
TRM single $z$	1	71.9

Image Source: Jolicoeur-Martineau, A. (2025). Less is More: Recursive Reasoning with Tiny Networks. *arXiv:2510.04871*.





```

def latent_recursion(x, y, z, n=6):
    for i in range(n): # latent reasoning
        z = net(x, y, z)
    y = net(y, z) # refine output answer
    return y, z

def deep_recursion(x, y, z, n=6, T=3):
    # recursing T-1 times to improve y and z (no gradients needed)
    with torch.no_grad():
        for j in range(T-1):
            y, z = latent_recursion(x, y, z, n)
    # recursing once to improve y and z
    y, z = latent_recursion(x, y, z, n)
    return (y.detach(), z.detach()), output_head(y), Q_head(y)

# Deep Supervision
for x_input, y_true in train_dataloader:
    y, z = y_init, z_init
    for step in range(N_supervision):
        x = input_embedding(x_input)
        (y, z), y_hat, q_hat = deep_recursion(x, y, z)
        loss = softmax_cross_entropy(y_hat, y_true)
        loss += binary_cross_entropy(q_hat, (y_hat == y_true))
        loss.backward()
        opt.step()
        opt.zero_grad()
        if q_hat > 0: # early-stopping
            break

```

# Model Design

## Tiny Recursive Model

- Use single network
- Less is more: 2-layer-transformer
- No attention on Sudoku
- No ACT (Q learning)
- Iteration steps

Table 1. Ablation of TRM on Sudoku-Extreme comparing % Test accuracy, effective depth per supervision step ( $T(n + 1)n_{layers}$ ), number of Forward Passes (NFP) per optimization step, and number of parameters

Method	Acc (%)	Depth	NFP	# Params
HRM	55.0	24	2	27M
TRM ( $T = 3, n = 6$ )	87.4	42	1	5M
w/ ACT	86.1	42	2	5M
w/ separate $f_H, f_L$	82.4	42	1	10M
no EMA	79.9	42	1	5M
w/ 4-layers, $n = 3$	79.5	48	1	10M
w/ self-attention	74.7	42	1	7M
w/ $T = 2, n = 2$	73.7	12	1	5M
w/ 1-step gradient	56.5	42	1	5M

Table 3. % Test accuracy on Sudoku-Extreme dataset. HRM versus TRM matched at a similar effective depth per supervision step ( $T(n + 1)n_{layers}$ )

		HRM		TRM	
		$n = k, 4 \text{ layers}$		$n = 2k, 2 \text{ layers}$	
$k$	$T$	Depth	Acc (%)	Depth	Acc (%)
1	1	9	46.4	7	63.2
2	2	24	55.0	20	81.9
3	3	48	61.6	42	87.4
4	4	80	59.5	72	84.2
6	3	84	62.3	78	OOM
3	6	96	58.8	84	85.8
6	6	168	57.5	156	OOM



# Performance

## Tiny Recursive Model

Table 4. % Test accuracy on Puzzle Benchmarks (Sudoku-Extreme and Maze-Hard)

Method	# Params	Sudoku	Maze
Chain-of-thought, pretrained			
Deepseek R1	671B	0.0	0.0
Claude 3.7 8K	?	0.0	0.0
O3-mini-high	?	0.0	0.0
Direct prediction, small-sample training			
Direct pred	27M	0.0	0.0
HRM	27M	55.0	74.5
TRM-Att (Ours)	7M	74.7	85.3
TRM-MLP (Ours)	5M/19M <sup>1</sup>	87.4	0.0

Table 5. % Test accuracy on ARC-AGI Benchmarks (2 tries)

Method	# Params	ARC-1	ARC-2
Chain-of-thought, pretrained			
Deepseek R1	671B	15.8	1.3
Claude 3.7 16K	?	28.6	0.7
o3-mini-high	?	34.5	3.0
Gemini 2.5 Pro 32K	?	37.0	4.9
Grok-4-thinking	1.7T	66.7	16.0
Bespoke (Grok-4)	1.7T	79.6	29.4
Direct prediction, small-sample training			
Direct pred	27M	21.0	0.0
HRM	27M	40.3	5.0
TRM-Att (Ours)	7M	44.6	7.8
TRM-MLP (Ours)	19M	29.6	2.4

# Summary

## Tiny Recursive Model

- Compared with HRM
  - Simpler: no fixed-point theorem, no biological hierarchy, no extra halting pass.
  - Smaller: one shared network  $\rightarrow \approx 1/4$  the parameters of HRM.
  - More robust: better generalization on four reasoning benchmarks.

# Conclusion

- A single tiny network **recursively** refines its latent reasoning state to improve answers **step by step**—achieving strong generalization through thinking depth rather than network depth.
- Why recursion helps so much remains to be explained (overfitting?).
- A supervised learning method rather than generative model.

# Thanks!